

Advancing Generative AI: A Critical Framework for Perceptual Evaluation, Latent Space Governance, and Semiotic Interfaces

1. Executive Summary: Charting New Frontiers in Generative AI Research Through Advanced Prompting Frameworks

This report synthesizes a comprehensive "Deep Research PROMPT SET" designed to critically evaluate and advance the capabilities of generative Artificial Intelligence (AI). It delves into the foundational epistemology of perceptual metrics, explores methodologies for the governance of latent textural and material properties within generative models, examines the role of material prompting as a semiotic interface for nuanced visual synthesis, and proposes rigorous psychophysical validation techniques for generative outputs. The overarching framework presented herein advocates for a significant evolution beyond current paradigms, aiming to foster the development of more robust, interpretable, and perceptually aligned AI systems.

Key analyses reveal that prevalent perceptual metrics such as the Learned Perceptual Image Patch Similarity (LPIPS) and Deep Image Structure and Texture Similarity (DISTS) exhibit considerable epistemic limitations and are susceptible to critical misalignment issues. This necessitates a strategic shift towards more sophisticated metrics, potentially embodied by a suite combining Multimodal Task-Aligned Similarity (MTAS), Structural Similarity Index (SSI), and Latent Consistency Model (LCM)-based approaches, alongside the exploration of non-score-based latent alignment strategies.

The conceptualization of an "Organic Realism Prior" (ORP) emerges as a potent framework for governing the synthesis of high-fidelity biological textures. Implementing such a prior requires a delicate balance between achieving verisimilitude and ensuring pedagogical utility, a challenge that demands innovative governance protocols.

Furthermore, treating material prompting not merely as instruction but as a sophisticated semiotic interface offers the potential to unlock more nuanced and affectively resonant visual synthesis. Realizing this potential, however, hinges on the development of standardized syntactical structures for material prompts and a deeper empirical understanding of the comparative efficacy of metaphor-based versus exemplar-based anchoring techniques.

The validation of generative outputs, particularly concerning temporal fidelity and perceived realism under dynamic conditions, calls for the integration of recursive perception simulations incorporating human-in-the-loop psychophysics. Such methods are crucial for assessing whether generated visuals maintain their perceived integrity across varying contexts.

Finally, the "Prompt Upgrade Module" embedded within the research set introduces critical meta-level inquiries. These challenge prevailing assumptions regarding visual realism, advocate for the constructive recontextualization of "authenticity gaps," and propose a re-conceptualization of semantic failures not as errors but as potentially valuable inputs for the generative process.

Based on these findings, this report recommends prioritizing research into unified and generalizable perceptual metrics, developing robust and adaptable latent space governance protocols for complex material properties, standardizing material prompting interfaces to enhance control and interoperability, and establishing rigorous, domain-aware psychophysical validation pipelines. These endeavors are essential for steering the next generation of AI towards producing outputs that are not only technically proficient but also meaningfully aligned with human

perception and diverse application requirements.

2. Prompt 1 Analysis: Re-evaluating Perceptual Metrics – Epistemology, Overfitting, and the Quest for Latent Alignment

The critical assessment of generative models hinges on the metrics used for their evaluation. This section, adopting the perspective of an advanced AI model evaluator with expertise in psychophysics, visual neuroscience, and generative model performance analysis, undertakes a critical and interrogative examination of existing and proposed perceptual metrics.

2.1. Comparative Failure Analysis: LPIPS, DISTS vs. MTAS/SSI/LCM

A thorough understanding of current perceptual metrics necessitates a comparative analysis of their methodologies, strengths, and, crucially, their documented failure modes.

LPIPS (Learned Perceptual Image Patch Similarity) utilizes deep neural network activations, typically from architectures like AlexNet, VGG, or SqueezeNet, as a foundation for assessing perceptual similarity between images.¹ The process often involves a linear calibration layer and computes differences between feature embeddings extracted from image patches, subsequently aggregating these differences into a similarity score.² Its widespread adoption stems from a generally better correlation with human perception compared to traditional metrics such as Peak Signal-to-Noise Ratio (PSNR) or the Structural Similarity Index (SSIM).² LPIPS can also be employed as a "perceptual loss" function during model training.¹

Despite its utility, LPIPS is predicated on the assumption of perfect image alignment and exhibits marked sensitivity to small, often imperceptible, spatial shifts between the compared images.² This sensitivity is a significant limitation, as even a one-pixel translation can drastically alter classifier predictions, a characteristic LPIPS inherits from its backbone networks.² Normalization practices also introduce variability; LPIPS typically requires input images to be scaled to a fixed intensity range (e.g., [-1, +1]), and choices in normalization can significantly affect scores, particularly when images possess differing intrinsic intensity ranges or distributions.⁵ This implies a risk of inconsistent evaluations if preprocessing is not rigorously standardized. Furthermore, while an improvement over SSIM, LPIPS can encounter difficulties with texture-rich regions. When used as an optimization target, it has been observed to yield low-quality results or even fail to fully reconstruct images in certain scenarios.⁶ The features LPIPS extracts may also carry biases from the pre-training tasks of its backbone networks, potentially leading to evaluations that do not consistently align with human perception, a concern also raised for metrics like FID that often rely on Inception-V3.⁷

DISTS (Deep Image Structure and Texture Similarity) aims to provide a more holistic image quality assessment by unifying structure and texture similarity, with an explicit design goal of tolerance to texture resampling.⁹ It employs a Convolutional Neural Network (CNN) to transform images into multi-scale, overcomplete representations. The metric then combines correlations of spatial averages of feature maps (termed "texture similarity") with correlations of the feature maps themselves ("structure similarity").⁹ The parameters of DISTS are jointly optimized to align with human ratings of image quality while simultaneously minimizing reported distances between subimages extracted from the same texture exemplars.⁹ This design allows DISTS to better handle texture variance and mild geometric transformations compared to LPIPS¹³, making it less sensitive to texture resampling, a common aspect of human perceptual tolerance.⁹

However, DISTS is not without its own set of issues. Optimization processes utilizing DISTS as a loss function have been reported to introduce undesirable visual artifacts, suggesting that the metric may not always adequately maintain image fidelity.⁶ Common with other learning-based Full-Reference Image Quality Assessment (FR-IQA) models, DISTS might place an overemphasis on quality separability, potentially neglecting the compactness of representations for images of very similar quality. This could arise from projecting and comparing images within a suboptimal subspace.⁶ While the original DISTS paper addresses texture resampling tolerance, separate research efforts have focused on building upon LPIPS to specifically tackle small spatial shift sensitivity by investigating techniques like anti-aliasing, pooling modifications, and striding adjustments.² This suggests that the original DISTS formulation may not have entirely resolved shift sensitivity, or that these are complementary challenges. The development of "Adaptive DISTS (A-DISTS)"¹⁴, which refines DISTS by adaptively localizing structure and texture regions and weighting similarities accordingly, further implies that the initial DISTS had scope for improvement in handling diverse image content. Similar to LPIPS, DISTS requires a normalized input range (e.g.,), making it susceptible to comparable normalization-related concerns.⁵

The proposed **MTAS/SSI/LCM suite** represents a conceptual shift towards more robust and multifaceted evaluation:

- **MTAS (Multimodal Task-Aligned Similarity):** While one set of information points to MTA in the context of Bird's Eye View (BEV) perception and captioning, focusing on aligning visual features with linguistic data¹⁵, a more pertinent interpretation for general perceptual evaluation aligns with the UniSim-Bench and UniSim models.¹⁶ These aim to create a *unified multi-task perceptual similarity metric*. The UniSim approach involves fine-tuning Vision-Language Models (VLMs) across a spectrum of perceptual tasks, demonstrating superior generalization capabilities compared to models trained on narrow, specific tasks.¹⁶ The underlying epistemic assumption is that human perception is a multifaceted, interconnected system; consequently, a robust metric should exhibit strong performance and generalization across these diverse perceptual functions rather than excelling in a singular, isolated domain.
- **SSI (Structural Similarity Index Metric):** SSIM and its multi-scale variant, MS-SSIM, evaluate similarity based on three key image characteristics: luminance, contrast, and structure.⁴ It generally offers better alignment with human perception of structural alterations than PSNR.¹⁹ However, SSIM remains sensitive to misalignments, albeit less so than PSNR²⁰, and can be overly reactive to texture resampling.⁹ Its mathematical nature as not being a true distance metric can limit its direct applicability in certain optimization contexts without transformation.¹⁹ Furthermore, SSIM scores are notably influenced by normalization procedures and the choice of data range parameters.⁵ The core epistemic assumption of SSIM is that perceived similarity is primarily determined by shared structural information, along with luminance and contrast characteristics.
- **LCM (Latent Consistency Models):** LCMs are primarily a technique for accelerating image generation by distilling pre-trained Latent Diffusion Models (LDMs), enabling few-step synthesis.²² Their relevance to evaluation stems from their intrinsic consistency principles and the potential for integrating perceptual loss during their training.²² The work on "Latent Guidance in Diffusion Models for Perceptual Evaluations" introduces Perceptual Manifold Guidance (PMG), which utilizes pretrained LDMs for No-Reference Image Quality Assessment (NR-IQA).²³ This approach hypothesizes that LDMs implicitly learn and exhibit perceptually consistent local regions within their data manifolds. PMG leverages "diffusion hyperfeatures"—multi-scale, multi-timestep features extracted from the U-Net of the LDM—which have been shown to correlate with human perception. As an evaluation concept, LCM-based methods could harness the rich, learned data manifolds of powerful generative models for perceptual assessment, potentially offering greater robustness and generalizability due to the extensive training of the foundation LDMs. However, the quality of such an evaluation would be intrinsically linked to the teacher LDM and the distillation process. A key

research challenge lies in ensuring that the "consistency" measured by these models aligns broadly with human perceptual consistency across diverse tasks. The epistemic assumption here is that the latent spaces of highly capable generative models, particularly diffusion models, encode perceptually relevant manifolds, and consistency within these learned manifolds can serve as a reliable proxy for perceptual quality or similarity.

The trajectory of perceptual metric development, moving from feature-distance measures like LPIPS and DISTS towards concepts such as the task-generalizable MTAS (UniSim) and model-intrinsic LCM-based evaluation, signals a notable *epistemic drift*. This shift is a direct response to the limitations of existing metrics and a critique of "metric solutionism," where the pursuit of high scores on narrow metrics has led to models that perform well by those specific measures but may fail perceptually in broader, more complex contexts (e.g., demonstrating overfitting or a lack of robustness to common image variations like small shifts). The very existence of dedicated research to enhance LPIPS's shift tolerance² or to adapt DISTS for varied content¹⁴ underscores the ongoing effort to address the persistent shortcomings of these established metrics.

Metric	Core Principle	Backbone (Example)	Training Data (Example)	Key Strengths	Documented Weaknesses/Failures	Core Epistemic Assumption
LPIPS	Learned deep features distance	AlexNet, VGG ¹	BAPPS dataset ¹	Good general perceptual correlation ² ; Perceptual loss ¹	Shift-sensitive ² ; Normalization issues ⁵ ; Texture region artifacts ⁶ ; Potential backbone biases ⁷	Deep classification features are a proxy for human perceptual features. ²
DISTS	Structure & texture feature correlation; Texture resampling tolerance	CNN (VGG-like) ⁹	IQA databases ⁹	Handles texture variance & mild geometric transforms ¹³	Artifact introduction with optimization ⁶ ; Overemphasis on separability ⁶ ; Original version may have shift sensitivity ³ ; Normalization dependent ⁵	Statistical properties of multi-scale deep features (structure & texture) equate to perceived quality. ⁹
SSIM	Luminance, contrast, structure comparison	N/A	N/A	Simple, interpretable structural comparison ¹⁹	Sensitive to misalignments & texture resampling ⁹ ; Not a true distance metric ¹⁹ ;	Local structural integrity, luminance, and contrast are primary determinants of perceived

					Normalization/ data range dependent ⁵	similarity. ¹⁹
MTAS (UniSim Concept)	Unified multi-task VLM fine-tuning for general perceptual similarity	CLIP, other VLMs ¹⁶	UniSim-Bench ¹⁶	Generalization across diverse perceptual tasks ¹⁶	Performance depends on VLM capabilities and breadth of fine-tuning tasks; Newer concept, fewer documented large-scale failure analyses specific to this unified approach compared to older metrics.	Human perception is a unified system; metrics should generalize across its facets. VLM embeddings capture rich perceptual info. ¹⁶
LCM (LGDM Concept)	Diffusion hyperfeatures from LDM manifold for No-Reference IQA	Pretrained LDM ²³	LDM training data	Leverages powerful LDM priors; Potential for high robustness ²³	Quality depends on teacher LDM & distillation; Ensuring alignment with broad human perception is an active research area. ²²	Learned manifolds of powerful generative models (LDMs) encode perceptually relevant information. ²³

Table 1: Comparative Analysis of Perceptual Metrics

2.2. Dissecting Epistemic Assumptions Baked into Each Metric

Every perceptual metric, by its design and training, implicitly embeds assumptions about the nature of human visual perception and how it can be computationally modeled.

LPIPS and DISTS share several core assumptions. Firstly, they presume that deep features extracted from networks—originally trained for large-scale image classification (like ImageNet, often used for LPIPS backbones) or tailored for specific IQA objectives—can serve as a valid proxy for the feature space utilized by human perception.² This "unreasonable effectiveness" of deep features²⁸, however, is increasingly scrutinized, with findings indicating that metrics reliant on specific architectures like Inception-V3 do not always correlate well with human judgments, particularly in the context of diffusion models.⁷ The specific choice of layers and their weighting within LPIPS¹ introduces a further layer of engineered assumptions regarding the relative importance of different features. Conventional IQA metrics, many of which are deep learning-based, have also faced critique regarding their alignment with both downstream DNN performance and true human perception, suggesting a persistent gap.²⁹ Secondly, these metrics assume that perceptual difference can be meaningfully quantified through a (weighted) linear combination of feature distances (as in LPIPS) or correlations (as in DISTS).¹ Human perception, however, is known to be highly non-linear and context-dependent. While these metrics offer improved correlation over simpler alternatives, this inherent simplification may fail to capture complex perceptual interactions. The "perception bias"

noted in DISTS, potentially arising from comparisons within an "inappropriate subspace"⁶, points to this limitation. Thirdly, as full-reference metrics, LPIPS and DISTS typically assume the availability of a perfectly aligned, distortion-free reference image, and that point-wise or patch-wise comparisons are semantically meaningful.² This assumption is challenged by even slight misalignments² and is less applicable to the evaluation of generative models where a perfect reference often does not exist, or where textural and stylistic variations are not necessarily defects but desired characteristics. DISTS's explicit attempt to tolerate texture resampling⁹ is an acknowledgment of this particular constraint.

SSIM operates on the assumption that the human visual system is highly adapted for extracting structural information from a scene. Consequently, changes in structure, luminance, and contrast are considered the primary determinants of perceived image quality.⁴ While this holds true to a degree, it tends to underweight the contributions of texture, color fidelity beyond simple luminance, and more complex semantic aspects of image quality. It has been noted that the HVS prioritizes structural information for lower-quality images but is more sensitive to point-wise variations in high-quality images.³⁰ SSIM also assumes that local statistics computed within a window are sufficient to capture these key attributes. This approach, however, does not adequately capture global context and long-range dependencies, which are crucial for human perception, though MS-SSIM¹⁹ attempts to partially address this by incorporating multiple scales.

The conceptual **MTAS (UniSim-like)** framework introduces different epistemic assumptions. It posits that various perceptual tasks—such as image-to-image similarity, text-image alignment, and image quality assessment—are merely different facets of a unified human perception system.¹⁶ A significant implication of this is that a metric trained across these diverse tasks should exhibit greater robustness and generalizability, marking a substantial epistemic departure from specialized, single-task metrics. Furthermore, this approach assumes that embeddings derived from large vision-language models (VLMs) capture rich, multimodally-aligned semantic and perceptual information.¹⁶ This moves beyond purely visual features to incorporate linguistic and conceptual understanding, potentially leading to better alignment with higher-level human judgments.

LCM-based evaluation concepts (exemplified by LGDM) are founded on the manifold hypothesis: that real-world data lies on a lower-dimensional manifold, and that generative models, especially LDMs, learn to represent these manifolds effectively.²³ Perceptual consistency, therefore, can be identified and measured within these learned manifolds. This implies that evaluation can become intrinsic to the model's own understanding of the data distribution, rather than relying on external, potentially mismatched feature extractors. A related assumption is that "diffusion hyperfeatures"—multi-scale, multi-timestep features extracted from the denoising U-Net of an LDM—are perceptually salient.²³ The underlying idea is that the denoising process itself, across various noise levels and feature scales, captures information that is relevant to the human perception of image quality.

Philosophical considerations further illuminate the epistemic underpinnings of these metrics. Visual experience provides non-inferential justification for beliefs about the world, distinguishing between awareness of specific particulars and properties, and awareness of structural features of experience, such as the visual field being perceived as bounded.³¹ This distinction is pertinent to how metrics might focus on "properties" (e.g., LPIPS features) versus "structure" (e.g., SSIM), and whether they adequately capture these holistic "structural features" of perception. The inherent understanding of limits and context in human perception, as suggested by the "boundedness" concept, might be something that current metrics fail to model. Epistemological assumptions regarding the nature of reality and the relationship between the observer and the observed (the knower and the known) fundamentally shape research methodologies and the interpretation of findings.³² Perceptual metrics implicitly adopt an objectivist epistemology, treating perceptual quality as an objective, measurable property of an image, with the "truth" of this quality best discovered via carefully engineered feature extractors and distance

functions. This often overlooks the subjective, constructed, and context-dependent nature of human perception. Furthermore, the concepts of epistemic uncertainty (uncertainty in the model's parameters or predictions) and aleatoric uncertainty (inherent noise or ambiguity in the data or task) are crucial.³³ Most perceptual metrics provide a single, deterministic score, thereby ignoring the uncertainty associated with their own predictions (a form of epistemic uncertainty of the metric itself) or the inherent variability and ambiguity present in human perceptual judgments (aleatoric uncertainty). This represents a critical epistemic blind spot.

Each metric, through its specific architecture and training data, implicitly defines an ontology of "what matters" in visual perception. LPIPS, for instance, prioritizes hierarchical features learned for image classification tasks. SSIM elevates local structural integrity to a primary concern. The UniSim/MTAS concept proposes a broader, more encompassing ontology that includes cross-modal understanding and multi-task generalization. LCM-based methods, in turn, suggest an ontology derived from the generative process itself. The documented failures of these metrics to consistently and robustly align with human judgment⁷ strongly suggest that these implicit, engineered ontologies of perception are either incomplete or fundamentally flawed. They may capture certain aspects of visual properties but miss the holistic, structural, and context-aware nature that characterizes human visual experience. The call for unified metrics¹⁶ is, in essence, a call for a more comprehensive and accurate perceptual ontology.

2.3. Model Overfitting to Flawed Perceptual Metrics

The reliance on imperfect perceptual metrics for training and evaluating generative models can lead to a phenomenon akin to overfitting, where models become adept at satisfying the metric's criteria without achieving genuine perceptual quality.

Evidence suggests that generative models can indeed overfit to their training data, leading to the reproduction of existing examples rather than the generation of truly novel content.³⁵ Standard evaluation metrics like Fréchet Inception Distance (FID) have been shown to be insensitive to this type of overfitting; a "copycat" model that merely memorizes and reproduces training samples can achieve state-of-the-art FID scores.³⁶ The optimization pressure to minimize perceptual losses, such as LPIPS, can inadvertently encourage "digital forgery" or the memorization of training data.³⁶

While not "overfitting" to the metric in the exact same sense as overfitting to training data, a generative model optimized using a flawed perceptual loss function (e.g., an LPIPS variant that is highly sensitive to shifts) will inevitably learn to produce images that are deemed "good" according to that specific, flawed criterion. These images, however, may be perceptually poor in other aspects that the metric fails to capture or misrepresents. For instance, a model might learn to generate overly smooth images if its perceptual loss function inaccurately penalizes high-frequency details that it misinterprets as noise, perhaps due to minor misalignments. The unreliability of feature-distance metrics like FID for assessing generative models, especially concerning their utility for downstream tasks or their true perceptual fidelity, is increasingly recognized.⁸ Models optimized solely against these metrics may not generalize well to real-world perceptual expectations or produce outputs that are genuinely useful or realistic. Even in Score-based Generative Models (SGMs), achieving precise empirical score matching does not guarantee the generation of innovative and diverse samples, suggesting a form of "overfitting" to the score objective itself, thereby missing broader creative and qualitative aspects of generation.³⁹

The implications of this metric-induced overfitting are significant. Models effectively learn to "game" the metric, producing outputs that achieve high scores but lack true perceptual quality, diversity, or novelty.³⁵ This creates a skewed perception of progress within the field of generative modeling, where improvements in reported scores do not necessarily translate to genuine enhancements in user-perceived quality or practical utility.⁷ Consequently, valuable research and computational resources may be expended on optimizing for flawed objectives. In

specialized domains such as biological or medical imaging, this could result in the generation of images that appear favorable to a metric but lack essential diagnostic utility or, more dangerously, introduce subtle, misleading artifacts that could confound interpretation.³⁷

Models that are over-optimized based on current perceptual metrics might inadvertently enter a "perceptual uncanny valley." They may appear close to realistic according to certain quantitative measures but "feel wrong" or subtly unsettling to human observers. This occurs because these models conform to the explicit criteria of the flawed metrics but deviate from the true, more complex perceptual manifolds that these metrics fail to fully capture. This "slight wrongness," an idea also surfaced in the "Failure-Informed Inversion Prompt," could be a direct consequence of models learning to exploit the limitations and blind spots of incomplete or flawed perceptual models. The consistent findings from human evaluation studies that highlight mismatches between FID scores and human judgments of realism⁷ lend support to the notion that current metrics might be leading generative models astray, towards regions of the generation space that are metrically optimal but perceptually suboptimal or even disconcerting.

2.4. Proposing Non-Score-Based Latent Alignment Strategies

The demonstrable limitations of score-based perceptual metrics⁷ underscore the urgent need for alternative, non-score-based strategies for evaluating and guiding generative models, particularly focusing on latent alignment.

Human-in-the-Loop (HIL) evaluation stands as a cornerstone. Psychophysical experiments, as envisioned in Prompt 4 of the research set, offer a direct way to assess perceptual qualities.⁷ The HYPE benchmark, for example, directly measures human perception of realism in generated samples.⁴² Another avenue involves model-based metrics that employ a "judge model," often a sophisticated LLM or VLM like Gemini, to assess outputs against qualitative criteria such as fluency, relevance, and correctness.⁴⁴ While these may produce scores, the scores are derived from a model's qualitative assessment rather than direct feature-distance computations. Pairwise comparison tasks, performed either by human evaluators or by judge models, also offer valuable qualitative insights.⁴⁴

Task-based evaluation provides a pragmatic approach by assessing generative models based on their utility in downstream applications.³⁷ For instance, the effectiveness of synthetic data in augmenting the training set for a classification or segmentation task can serve as a measure of its quality and relevance.

Direct **latent space probes and analysis** offer a more intrinsic evaluation method. This involves investigating the geometric and topological properties of the learned latent space. Key questions include whether semantically similar inputs map to proximal points in the latent space and whether meaningful and smooth interpolations can be performed between latent codes.⁴⁶ A critical aspect is the exploration of **disentanglement of latent factors**.⁴⁹ A well-aligned generative model should ideally possess a latent space where distinct perceptual attributes are disentangled and independently controllable. Changes made to specific dimensions in such a latent space should correspond to predictable and isolated modifications in the generated output. The **Perceptual Manifold Guidance (PMG)** method, emerging from the LGDM framework²³, aligns with this philosophy. PMG uses pretrained LDMs and perceptual quality features to guide on-manifold sampling, yielding feature maps that correlate with human perception for NR-IQA. While PMG can produce a score, its underlying principle is about aligning generation with the *learned perceptual manifold* of a powerful generative model—a fundamentally qualitative alignment goal. The "diffusion hyperfeatures" themselves could be used for non-score-based characterization of perceptual properties.

Causal analysis of the generative process aims to move beyond mere correlation by understanding how specific changes in input prompts or latent codes causally influence particular perceptual attributes in the output. This fosters a deeper understanding of the model's internal workings.

For specialized domains, such as medical imaging³⁷ or scientific visualization⁵⁵, **qualitative analysis by domain experts** remains indispensable. Their nuanced understanding of domain-specific requirements for realism and utility cannot be easily encapsulated by automated scores.

The "Reflexivity-as-End Layer" lens, applied here, implies that the evaluation process itself must be subject to continuous scrutiny and iterative refinement. Non-score-based methods, especially those involving human judgment or qualitative expert review, inherently embody this reflexivity to a greater extent than fixed, score-based metrics.

A particularly promising non-score-based strategy involves a form of "latent space cartography." This approach focuses on mapping the latent space of generative models to identify regions that correspond to high perceptual quality, semantic coherence, and desired attributes, without necessarily collapsing this complex understanding into a single numerical score. It requires a deep understanding of the geometry, topology, and semantic disentanglement of the latent space.⁴⁶ Methodologies like PMG²³, by identifying "perceptually consistent local regions within the data manifold," represent initial steps in this direction. Current metrics are demonstrably insufficient⁷, necessitating alternatives.³⁹ Given that the latent space is the origin of generation⁴⁶, a well-organized and understood latent space—characterized by properties like disentanglement, continuity, and completeness⁴⁶—would allow "good" generation to be defined by *location within or trajectories through* this space. This is not about achieving a singular score but about understanding the *internal structure* of perceptual quality as represented within the model itself. PMG's use of "diffusion hyperfeatures"²⁴ to correlate with human IQA judgments exemplifies probing the latent and feature spaces for perceptually relevant signals, a form of qualitative alignment that moves beyond simple scoring.

3. Prompt 2 Analysis: Governing Latent Texture – The "Organic Realism Prior" and Adaptive Material Guidance

This section explores the governance of latent textural properties, particularly for high-resolution biological textures, from the perspective of a hybrid prompt engineer and AI texture theorist. The focus is on modeling an "Organic Realism Prior" (ORP), balancing realism with pedagogical needs, and designing adaptive guidance mechanisms.

3.1. Modeling the "Organic Realism Prior" (ORP) – A Latent Construct

The "Organic Realism Prior" (ORP) is conceptualized as a learned distribution or a set of constraints operating within the latent space of a generative model, such as a diffusion-GAN fusion architecture. Its purpose is to steer the synthesis process towards textures that exhibit the multifaceted characteristics of biological "realism." This extends beyond mere photorealism to encompass properties like naturalistic variation, specific microstructural details, and adherence to underlying biological principles, even when these principles are represented abstractly.

The latent space, being a compressed representation that preserves essential features⁴⁶, would, under an ORP, prioritize features pertinent to organic qualities. For effective governance by the ORP, this latent space should ideally exhibit continuity, meaning that proximal points in the space yield similar organic textures, and completeness, ensuring that any point sampled from this space decodes into a meaningful organic texture.⁴⁶ The ORP could be enforced by actively shaping the latent space distribution itself (akin to how Variational

Autoencoders might force a Gaussian distribution ⁴⁶) or through the learned dynamics of a diffusion model operating within this constrained space.⁴⁷

Inspiration for structuring such a prior can be drawn from frameworks like OrgMol-Design, which is used for generating complex organic molecules.⁵⁸ This system employs a score-based generative model incorporating a "fragment prior" for coarse-grained scaffold generation, complemented by a "chemical-rule-aware scoring model" for fine-grained molecular structure design. This suggests that the ORP could also be multi-scalar in nature: a "macro-ORP" could govern overall tissue or organ morphology, while a "micro-ORP" could manage cellular or sub-cellular texture details, potentially guided by learned biological rules or statistical properties derived from real biological textures. The concept of a "realism index" for points within a latent space, which can be constructed from a prior density or an FID-like approach to ensure that interpolations remain close to the data manifold ⁴⁸, offers another mechanism. An ORP could be operationalized as a region or manifold within the latent space characterized by a high realism index with respect to organic textures.

A governance system for the ORP would likely involve several key components: an encoder to map example organic textures into the ORP-conditioned latent space; a prior model that explicitly defines or constrains the ORP distribution; a decoder or generator to synthesize textures from points within this ORP-governed space; and a set of steering protocols, including advanced prompts like meta-control prompts, to modulate sampling from or traversal within this specialized latent space.

The ORP must be designed to support a "Typological Drift" from schematic (abstract, primarily pedagogical) representations to immersive (highly realistic, detailed) ones. This implies that the latent space governed by the ORP should not be monolithic but should allow for controllable variation along this axis of abstraction. Furthermore, the "Latent Space Ontology" lens suggests that the dimensions or regions of this ORP-governed latent space should ideally correspond to meaningful, perhaps disentangled, aspects of organic texture, such as cell density, fiber orientation, or vascularity.

The ORP should not be conceived as a static, predefined prior but rather as a dynamic, learnable constraint system. It could potentially integrate elements of physical or biological simulation, or rule-based systems ⁵⁸, directly into the definition of the latent space or into the generative process itself. The latent space, as described in ⁴⁶, preserves "essential features"; for the ORP, these are the defining characteristics of organic realism. The OrgMol-Design system's use of "fragment priors" and "chemical-rule-aware scoring" ⁵⁸ points towards an ORP that is not merely about matching a dataset of images but about encoding underlying biological principles. If the ORP defines a manifold of organic realism, as suggested by the "realism index" concept ⁴⁸, this manifold needs to be more than just a cluster of data points; it must embody the "rules" governing organic textures. Consequently, the governance system involves not just sampling from a fixed distribution but interacting with a system that understands, or can be guided by, these rules, rendering it inherently dynamic and learnable.

3.2. Heuristics for Balancing Realism Fidelity and Pedagogical Abstraction

Achieving a balance between high realism fidelity and effective pedagogical abstraction is crucial for generative models producing biological textures. Realism fidelity strives for high verisimilitude with actual biological textures, capturing fine details, natural stochasticity, and the complex interplay of light and material.⁵¹ This is vital for immersive applications and advanced research simulations. Conversely, pedagogical abstraction aims for clarity, often simplifying or stylizing textures to highlight key structures or processes, thereby improving understanding in educational contexts. This might involve reframing the "slight wrongness" (identified in the Failure-Informed Inversion Prompt) as an intentional, purposeful abstraction.

Several heuristics can be employed to navigate this balance:

- **Controllable Disentanglement:** If the latent space successfully disentangles "realism factors" (e.g., level of micro-detail, complexity of noise) from "structural factors" (e.g., cell type, tissue organization), users can navigate this space to tune the desired balance.⁴⁹ Prompts could then target specific levels of abstraction for certain features while maintaining realism for others.
- **Multi-Scale Representation:** Inspired by multi-scale feature analysis in metrics like DISTS⁶¹ and multi-stage generation in complex molecule design⁵⁸, textures could be generated with a coarse, abstracted base layer, with finer, realistic details progressively added. The amount of detail incorporated at each scale could be a controllable parameter.
- **Latent-Space or Post-Generation NPR:** Techniques from Non-Photorealistic Rendering (NPR) could be applied, either by manipulating latent representations or as a post-processing step, to abstract realistic textures for pedagogical aims.
- **Exemplar-Guided Abstraction:** The system could be guided by providing both a "realistic" exemplar and an "abstracted" exemplar, allowing it to learn or interpolate the desired level of abstraction.
- **Semantic Control Prompts:** Prompts could be designed to explicitly request different levels of detail or emphasis, e.g., "Show typical cell structure, emphasizing mitochondria" (pedagogical focus) versus "Hyperrealistic rendering of endoplasmic reticulum surface texture" (realism focus).
- **Loss Function Weighting:** If the model is co-trained to satisfy objectives for both realism and abstraction (perhaps using different discriminators or guidance signals), the balance can be dynamically tuned by adjusting the relative weights of the respective loss terms.

The "Semiotic Control vs. Interpretability Tension" lens is particularly relevant here: highly realistic textures might be information-rich but semiotically dense and difficult to interpret for pedagogical purposes. Abstracted textures offer greater interpretability but may lack semiotic richness or fidelity. The proposed heuristics aim to manage this inherent tension. Furthermore, the "Cognitive Friction Auditor" lens reminds us that prompts designed to achieve a specific balance must remain intuitive. If attaining pedagogical abstraction requires overly complex or non-obvious prompting strategies, the cognitive friction for the user increases, diminishing usability.

Balancing realism and pedagogy is not merely a process of "dumbing down" a realistic texture. Instead, it should be viewed as a principled method of information subtraction (removing distracting or irrelevant detail) and information highlighting (emphasizing key features essential for understanding). This can be framed as a guided traversal within a well-structured latent space, moving from a "high-fidelity" region to a "schematic" or "pedagogically optimized" region along specific, meaningful semantic axes. The concept of disentangled representations⁴⁹, where factors like "level of detail" or "presence of specific organelles" are independently controllable, directly supports this notion. Such controlled movement within the latent space allows for selective emphasis and de-emphasis, akin to the process an expert human illustrator employs when creating a clear pedagogical diagram from a complex micrograph. The "Typological Drift (from schematic → immersive)" lens directly aligns with this idea of a controllable spectrum of representation.

3.3. "Meta-Control Prompts" for Adaptive Material-Level Guidance

"Meta-control prompts" are defined as higher-level prompts that function by generating or modifying other, more specific prompts which, in turn, directly guide the texture synthesis process. They introduce an additional layer of abstraction and control, enabling adaptive, context-aware, or user-driven refinement of material guidance without necessitating manual rewriting of detailed, low-level prompts.

The purpose of such meta-control prompts is to allow users to guide the generative process at a more conceptual level. For example:

- A user might issue a meta-control prompt like: "Generate a sequence of five prompts to visualize the progression of epithelial tissue healing, starting from acute inflammation and culminating in mature scar tissue. For each stage, adjust the underlying texture prompts to reflect characteristic changes in cell density, collagen fiber alignment and density, and degree of vascularity."
- Another example could be: "Given the base prompt P, which describes 'healthy liver tissue', create a modified prompt P' for 'cirrhotic liver tissue' by systematically increasing textual cues related to fibrosis texture and decreasing cues related to hepatocyte regularity and healthy coloration."
- For interactive refinement: "If user feedback indicates the current texture is 'too uniform' or 'lacks natural variation', generate a prompt variant that introduces a specified percentage increase in stochastic noise for material parameter X or incorporates subtle random variations in parameter Y."

The mechanism for implementing meta-control prompts could involve a Large Language Model (LLM) or a sophisticated rule-based system. This system would take the meta-control prompt and potentially existing base prompts (or even direct material parameters) as input, and output new or modified low-level prompts designed to achieve the higher-level goal.

This concept heavily relies on a robust "Prompt Infrastructure," ideally supported by "Prompt Ontologies." Such an infrastructure would require a structured, hierarchical ontology of prompt components (e.g., material categories, specific features, physiological states, visual characteristics like glossiness or roughness). This ontology defines the "vocabulary" and "grammar" for both low-level and meta-level prompts, enabling the meta-control system to parse, manipulate, and generate prompts in a meaningful way. This directly connects to the ideas in Prompt 3 concerning "Material Prompting as Semiotic Interface." A standardized syntax for material prompting (e.g., the proposed] from Prompt 3, or more generally discussed in ⁷⁶) would be essential for meta-control prompts to effectively deconstruct, modify, and reconstruct low-level prompts.

Meta-control prompts can facilitate an interactive, evolutionary approach to texture design. The user guides the generative process at a high conceptual level, and the system translates these high-level directives into specific material parameter changes or detailed textual cues via the automated generation or modification of low-level prompts. This allows for more intuitive personalization and efficient exploration of the complex latent space governed by the Organic Realism Prior. If a standardized syntax for material properties exists, an LLM or rule-based system can decompose user intent (e.g., "make the tissue appear more diseased") and reconstruct this into specific modifications within the low-level prompt syntax (e.g., "increase parameter X, decrease parameter Y"), based on learned associations or predefined mappings within the prompt ontology. This constitutes a form of adaptive guidance, enabling a more intuitive and less technically demanding interaction for users wishing to explore nuanced material variations.

4. Prompt 3 Analysis: Material Prompting as a Semiotic Interface – Encoding Cross-Modal Affective Synthesis

This section explores "Material Prompting" from the viewpoint of a semiotic systems designer, focusing on its role at the interface of cross-modal AI prompting and affective visual synthesis. The core idea is to treat material descriptions as rich semiotic codes rather than simple labels.

4.1. "Material Prompting" as a Semiotic Encoding Strategy

"Material Prompting" involves using textual descriptions of materials—such as "skin," "gelatin," "wet glass," "silk," or "rough stone"—not merely as categorical labels but as complex semiotic codes. These codes instruct the generative model to synthesize specific constellations of visual and implied haptic qualities. Each code carries a wealth of connotations related to texture, reflectance, translucency, pliability, temperature, and even affective

associations.

This concept is deeply intertwined with **semiotic theory**.⁶² Material prompting can be understood as an act of **intersemiotic translation**, converting verbal language (the prompt) into visual configurations (the generated image).⁶² In this translation, the AI model must establish meaningful correspondences between linguistic categories and a vast array of visual elements learned during its training.⁶² The material term in a prompt (e.g., "velvet") acts as a **signifier**. The **signified** is not a single, simple meaning but a complex bundle of visual properties (e.g., deep, saturated color, soft anisotropic sheen, characteristic highlight behavior) and affective associations (e.g., luxury, softness, warmth) that the model is tasked with rendering. Material prompts implicitly invoke what semioticians refer to as **plastic categories** (eidetic/form, chromatic/color, topological/spatial arrangement) and aspects of **visual enunciation** (how the material interacts with light, how its surface detail implies texture, how it is presented to the viewer).⁶² AI systems, through computational models, learn to understand and generate symbols; material terms are such symbols, acquiring their rich meanings through cultural convention, patterns which the AI discerns from its extensive training data.⁶⁴

The "Sensory Encoding from AI Prompting" lens highlights the objective: to encode not just the visual appearance but also implied sensory qualities. For instance, "wet glass" suggests coolness, smoothness, and transparency, while "rough stone" implies solidity, coarseness, and perhaps a certain temperature. This encoding can extend to more abstract affective qualities; "silk" might evoke feelings of elegance and smoothness, whereas "rust" could connote decay and age. The nature of material prompting is inherently **cross-modal**: the input is textual, the primary output is visual, but the underlying concept of the material often engages multiple sensory modalities (visual, tactile, and sometimes even associated auditory or olfactory memories).

The field of **text-to-material synthesis**⁶⁸ provides practical implementations of this semiotic interface. Systems like DreamMat, which generates PBR materials from text aiming to avoid baked-in shading effects⁷², and MatFuse, which utilizes diffusion models to create Spatially-Varying Bidirectional Reflectance Distribution Function (SVBRDF) maps with multiple conditioning sources (including text, sketches, and color palettes)⁷⁰, are concrete examples of material prompting functioning as a sophisticated semiotic encoding strategy.

4.2. Proposing a Standard Syntax/Interface Layer (e.g., I)

Currently, material prompting in generative AI often relies on ad-hoc natural language descriptions. The introduction of a standardized syntax or interface layer could offer more precise control, enhance modularity, and improve interoperability between different models or components within a generative pipeline. Such standardization is also a prerequisite for the effective functioning of meta-control prompts, as discussed in Prompt 2.

Inspiration for such a syntax can be drawn from general **prompt engineering principles**⁷⁶, which advocate for using structural patterns, separating instructions from context⁷⁶, and clearly defining roles, goals, context, and providing examples where appropriate.⁷⁷

A conceptual standardized syntax might include elements such as:

]

Where:

- <BaseMaterial> could be a primary material category (e.g., Glass, Skin, Metal, Fabric, Wood, Stone).
- <Property> could specify attributes like Roughness, Specularity, Translucency, Anisotropy, Emissiveness, RefractiveIndex.

- <Value> associated with a property could be numerical (e.g., a float between 0 and 1), categorical (e.g., Low, Medium, High), or descriptive (e.g., Specularity:Sharp_Focused, Roughness:Velvety_Soft, Translucency:Milky).
- <State> could describe modifications or conditions of the material (e.g., Wet, Dusty, Scratched, Tarnished, Molten, Frozen).

For instance, a prompt like] would instruct the model to generate wet glass with low roughness and high specularity. The example from the original prompt,], could be parsed or mapped within such a system to a more detailed internal representation like] (assuming "SemiGel" implies partial gelatin-like properties with some retained form, and "Specular" highlights a key desired visual characteristic).

This approach aligns with the "Prompt Ontology Structuring" lens. The proposed syntax would serve as the surface representation of a more comprehensive underlying ontology of materials, their inherent properties, their possible states, and the semantic relationships between these elements. Such an ontology would define the valid terms and their meanings, ensuring consistency. This structure promotes **modularity**, allowing complex material descriptions to be constructed from simpler, potentially reusable components. Different generative models or modules could implement support for various parts of this ontology, facilitating a more flexible and extensible system.

This structured approach is particularly relevant for the **controllable generation of PBR materials**.⁷⁰ PBR systems fundamentally rely on a set of parameters (e.g., albedo, metallic, roughness, normal maps, ambient occlusion) to define material appearance. A standardized prompt syntax could provide a user-friendly frontend that maps intuitive, semiotically rich descriptions to these complex PBR parameters. Systems like MatFuse⁷⁰ and MaterialMVP⁷⁴, which already aim for controllable PBR generation using multiple conditioning inputs, would benefit significantly from such a standard, as it would streamline the way users specify their intent.

The proposed syntax acts as a crucial bridge between human semiotic understanding and machine-parsable parameters. For human users, it offers a structured yet expressive way to articulate complex material concepts semiotically. For the generative models, it provides a more deterministic input that can be reliably mapped to underlying generation parameters (such as PBR values or specific coordinates/regions in a material-focused latent space). This structured communication reduces the inherent ambiguity often present in purely natural language prompts, thereby formalizing the "intersemiotic translation"⁶² by establishing a shared grammar between user and machine.

4.3. Failure Cases: Metaphor-Based Prompting vs. Exemplar-Based Anchoring

When prompting for materials, users might employ metaphorical language or provide visual exemplars. Both approaches have unique strengths and potential failure modes.

Metaphor-Based Prompting involves using figurative language to describe desired material qualities, for example, "skin like polished alabaster," "glass shattered like a broken promise," or "fabric flowing like water."

- **Potential Strengths:** Metaphors can evoke rich, affective qualities and complex combinations of properties that are often difficult to specify through literal description. They can tap into shared cultural understandings and allow for more poetic and nuanced expressions of intent.
- **Failure Cases**⁷⁸:
 - *Misinterpretation or Over-Literalness:* AI models may interpret metaphors too literally, leading to bizarre or unintended visual outputs (e.g., rendering an actual sculpture of alabaster instead of skin possessing alabaster-like smoothness and translucency). Research indicates that metaphors generated by models (e.g., textual metaphors from images) may not always capture the same semantic meaning as human-generated metaphors.⁷⁹

- *Ambiguity*: Metaphors are inherently open to multiple interpretations. The AI's chosen interpretation might not align with the user's intended meaning.⁸⁰
- *Lack of Grounding/Training Data Deficiencies*: If the AI model has not encountered sufficient examples linking the metaphorical source domain (e.g., alabaster) and the target domain (e.g., skin) in its training data, it may fail to bridge the conceptual gap and produce a meaningful synthesis.
- *Unintended Bypassing of Safety Filters ("Jailbreaking")*: Metaphor-based prompts have been identified as a potential vector for bypassing safety filters in text-to-image models.⁷⁸ This suggests that metaphors might operate in a semantic space that is less constrained by literal safety rules, implying a risk of generating unintended or harmful outputs if metaphors are not carefully managed or understood by the safety mechanisms.
- *Increased Cognitive Friction*: The "Cognitive Friction Lens" suggests that complex, obscure, or overly abstract metaphors can increase the cognitive load for the user in predicting the likely output and for the AI model in interpreting the prompt accurately.

Exemplar-Based Anchoring involves providing an example image or image patch of the desired material alongside, or instead of, a textual prompt.

- **Potential Strengths**: This method provides strong visual grounding for the desired appearance, significantly reducing ambiguity.⁸¹ It can allow for very fine-grained control over material properties if the exemplar is well-chosen and representative. Complex, subtle material characteristics that are difficult to articulate textually can often be effectively conveyed through an exemplar.
- **Failure Cases**:
 - *Over-reliance or Superficial Copying*: The model might simply copy the texture or appearance directly from the exemplar rather than understanding and transferring the underlying *material properties* to a new geometry or context.⁷⁶
 - *Poor Generalization*: If the provided exemplar is too specific (e.g., captured under unique lighting conditions, exhibiting particular wear-and-tear patterns not relevant to the target), the model may not generalize the core material properties effectively to new objects or different environmental conditions.
 - *Conflict between Exemplar and Textual Prompt*: If both a textual prompt and a visual exemplar are provided, conflicting information between them can confuse the model, leading to unpredictable or compromised results.
 - *Computational Overhead*: Processing and integrating image exemplars into the generation pipeline can be more computationally intensive than handling text-only prompts.

The "Semiotic Drift Lens" is relevant here: the meaning of a material term (e.g., "leather") can evolve or be interpreted differently. An exemplar anchors this meaning to a specific visual instance, thereby reducing semantic drift. Metaphors, being inherently less anchored and more abstract, are more susceptible to semiotic drift or misinterpretation based on the AI's internal "understanding" and training biases. For tasks like **semantic material editing**⁸³, exemplars or visually consistent references are often crucial. InterFaceGAN, for example, manipulates latent codes for semantic face editing, a concept that could be extended to materials if a well-structured "material latent space" is identified and can be navigated or conditioned by exemplars.⁸⁴ Systems like LatentEditor⁸⁵ and FocalDreamer⁸⁶ utilize prompts for 3D editing, where maintaining visual consistency with existing parts of an object (akin to exemplar guidance) is paramount.

Metaphor-based and exemplar-based methods are not necessarily mutually exclusive and could, in principle, be used synergistically (e.g., a prompt like "a surface with the *tactile feel* suggested by this [exemplar image of rough bark] but with the *luminous shimmer* of 'moonlight on water' [metaphor]"). However, an inherent tension exists:

exemplars tend to ground and concretize, while metaphors often abstract and poeticize. An advanced material prompting interface would need to intelligently manage this tension. A significant failure case could arise from an unresolved conflict between the literal visual information presented by an exemplar and the more abstract or relational qualities suggested by a metaphor, potentially leading to an incoherent, oscillating, or aesthetically unpleasing output. This points to the need for more sophisticated "semantic blending" or reconciliation mechanisms within the generative model when handling such combined inputs.

5. Prompt 4 Analysis: Recursive Perception Simulation – Validating Biological Visuals via Human-in-the-Loop Psychophysics

This section, adopting the role of a reflexive AI system simulator, focuses on designing and implementing protocols for validating high-fidelity biological visuals, particularly their temporal aspects, through rigorous human-in-the-loop psychophysical experimentation.

5.1. Simulation Protocol for Temporal Misalignment in Perception

A critical question in evaluating generative models is whether images rated highly on static perceptual metrics (e.g., achieving a high SSI score) maintain their perceived realism when subjected to dynamic contextual changes, such as variations in lighting or implied temporal states. This probes the robustness of static realism.

The **psychophysical methodology** for such a protocol would involve ⁷:

- **Stimuli Generation:**
 1. Generate a base high-fidelity biological image (e.g., a detailed micrograph of a cell, a cross-section of tissue) using a state-of-the-art generative model.
 2. Generate variants of this base image by systematically altering simulated lighting conditions (e.g., changing the angle, intensity, or color temperature of light sources to mimic time-of-day shifts or different microscopy illumination) or by introducing subtle shifts in the implied temporal state of the biological subject (e.g., progressing from an early to a mid-stage of apoptosis, showing slight organelle displacement or morphological change).
- **Experimental Task:** Employ a Two-Alternative Forced Choice (2AFC) task, a rating scale, or a realism judgment task.
 - Participants could be presented with the original image alongside a temporally or lighting-altered variant and asked: "Which of these two images appears more realistic under the specified new lighting condition?" or "Does this altered image look like a plausible continuation or a natural state variation of the original image?"
 - Alternatively, participants could rate the "realism" or "temporal consistency" of the altered images individually on a Likert scale (e.g., from "highly unrealistic" to "highly realistic").
- **Adaptation of HYPE:** The Human eYe Perceptual Evaluation (HYPE) methodology ⁴², which measures visual perception under adaptive time constraints or by quantifying human error rates in distinguishing real from fake images, could be adapted. In this context, an image deemed "temporally inconsistent" or "unrealistically lit" could be classified as "fake" or "less plausible."
- **Data Collection (Metrics):** Collect data on participants' accuracy (in 2AFC), realism ratings, reaction times, and confidence levels. The finding that no single existing computational metric strongly correlates with human evaluations of realism ⁷ underscores the critical importance of direct human testing for these nuanced perceptual judgments.

This protocol directly engages with "Temporal Fidelity Engineering" by testing whether the generative model possesses an underlying understanding that allows it to re-render a scene or biological state plausibly under new temporal (lighting, state evolution) conditions. The "Psychophysical Reflexivity" lens is also active, as the results from these human perception tests are intended to feed back into the iterative development and refinement of the generative models and their evaluation metrics. The work by Fan et al.⁸⁷ emphasizes modeling visual realism based on a deep understanding of human perception, which aligns with this reflexive goal. While this prompt focuses on static images that imply a dynamic context, insights can be drawn from research on **evaluating temporal coherence in video synthesis**.⁹³ For instance, methods developed for detecting temporal inconsistencies in deepfakes by analyzing facial patterns over time⁹⁵ could inspire ways to define and quantify "temporal misalignment" for static images that are part of an implied sequence or represent a dynamic biological process.

An image that is perceived as highly "real" in isolation might lose its perceived realism if it cannot be plausibly situated within a dynamically changing context (e.g., evolving lighting, progressing biological state). This suggests that static realism metrics, even sophisticated ones, are insufficient. Models need to learn a form of "conditional realism" that is robust to such dynamic alterations, demonstrating a deeper understanding of the underlying 3D scene, material properties, and biological processes that govern realistic appearance changes over time or under varying conditions.

5.2. Temporal Graph + Narrative-of-Thought (NoT) for Image Prompt Continuity Testing

To assess and improve the ability of generative models to produce coherent sequences of images representing evolving biological processes, a strategy combining Temporal Graphs and Narrative-of-Thought (NoT) prompting is proposed.

- **Temporal Graph:** This involves constructing a directed graph where each node represents a distinct image prompt corresponding to a discrete time point or stage in a biological process (e.g., T1: Quiescent_Macrophage → T2: Macrophage_Phagocytosing_Pathogen → T3: Macrophage_Presenting_Antigen). The directed edges in this graph signify the temporal transitions between these states.
- **Narrative-of-Thought (NoT):** For each transition (i.e., each edge in the temporal graph), a "Narrative-of-Thought" prompt is meticulously generated. This NoT prompt explicitly instructs the generative model on *how* to visually transform the image from state T_n to state T_{n+1} . It focuses on articulating the key visual changes that characterize the temporal progression and the underlying biological mechanisms.
 - For example, an NoT prompt for the transition T1 (Quiescent_Macrophage) → T2 (Macrophage_Phagocytosing_Pathogen) might be: "Transitioning from a quiescent state to active phagocytosis: show the macrophage extending pseudopods to engulf the pathogen, formation of a phagosome, and subtle changes in cell surface ruffling. Ensure pathogen is partially internalized."
- **Image Prompt Continuity Testing Protocol:**
 1. Generate an initial image I1 based on the prompt P1 (representing state T1).
 2. Utilize image I1 and the corresponding NoT_1_2 (the Narrative-of-Thought prompt for the transition from T1 to T2) as combined input to the generative model to produce image I2_NoT.
 3. Separately, generate an image I2_direct based solely on prompt P2 (representing state T2), without any explicit continuity guidance from I1 or the NoT.
 4. **Psychophysical Task:** Present participants with image I1 followed by a pair of images, I2_NoT and I2_direct (order randomized). Ask participants to judge: "Which of the two bottom images represents a more plausible and visually continuous next step after the top image, given the described biological process?" Alternatively, participants could rate the perceived visual continuity of each potential successor

image relative to I1.

5. This experimental design directly tests whether the NoT strategy significantly improves the perceived temporal coherence of the generated sequence compared to simply prompting for discrete, independent states.

This process inherently involves **recursive perception simulation**.⁹⁷ The output of one generation step (guided by an NoT prompt) can become the input or reference for the subsequent step in the sequence. Frameworks like Recursive Generative Emergence (RGE)⁹⁷, which discuss intelligence and complexity arising from recursive processes, provide a conceptual backdrop. The theory of Active Inference⁹⁸, which posits that the brain uses internal generative models to predict sensory input and then acts to make sensations conform to predictions, is relevant to how NoT attempts to guide the model to "predict" and render the next plausible visual state. However, it is also important to be mindful of potential failure modes: recursive training of generative AI models on their own outputs can sometimes lead to "model collapse," where the diversity and quality of generations degrade over iterations.⁹⁹ This is a critical factor to monitor in continuity testing, especially if generated outputs become too self-referential without sufficient grounding in external data or robust priors.

The "Failure-Informed Evaluation" lens is active here: if I2_NoT is consistently rated as worse than, or no better than, I2_direct in terms of continuity, it would indicate a failure in the NoT strategy itself or in the model's capacity to effectively interpret and follow such narrative guidance.

The NoT strategy provides a form of explicit cognitive scaffolding for the generative model, guiding its "visual reasoning" about temporal processes. This contrasts sharply with approaches that expect the model to infer complex temporal continuity solely from a series of prompts for discrete states. The success of an NoT-guided approach would indicate the significant value of more "programmatic," "story-like," or causally-informed prompting strategies when dealing with the generation of complex visual sequences that unfold over time. This is analogous to how humans often learn complex processes more effectively by understanding the intermediate steps and transformations, rather than just memorizing the start and end points.

5.3. Ambiguity Injection: Shifts in Perceptual Judgments with Slightly Wrong Temporal State Renderings

To further probe the robustness of perceived realism and temporal coherence, this part of the protocol involves injecting controlled ambiguities or errors into generated sequences.

- **Methodology for Ambiguity Injection:**

1. Generate a "correct" or canonical sequence of images representing a specific biological process (e.g., I1_correct, I2_correct, I3_correct, I4_correct).
2. For one or more states within this sequence, generate "slightly wrong" or ambiguous variants. Examples of such variants could include:
 - I2_early: An image for state T2 that incorrectly shows features more characteristic of an earlier stage (T1).
 - I2_late: An image for state T2 that prematurely shows features of a later stage (T3).
 - I2_exaggerated: An image that depicts the correct features for state T2 but renders them with unnatural exaggeration or intensity.
 - I2_missing_feature: An image for state T2 that omits a key characteristic feature expected at that stage.

- **Psychophysical Task Design:**

- Present participants with sequences containing these "slightly wrong" variants, for example: (I1_correct,

I2_early, I3_correct, I4_correct) or (I1_correct, I2_correct, I3_exaggerated, I4_correct).

- Ask participants to identify the image in the sequence that seems "out of place," "less plausible," or "temporally inconsistent" with the overall biological process being depicted.
- Participants could also be asked to rate the overall "smoothness," "correctness," or "biological plausibility" of the entire depicted sequence on a suitable scale.
- Advanced techniques like eye-tracking could be employed to determine if the "wrong" or ambiguous images attract more visual attention, longer fixation durations, or different gaze patterns compared to the correctly rendered images in the sequence.

While much of the existing literature on **ambiguity in text-to-image generation** focuses on resolving ambiguities present in the *input prompt*⁸⁰, this protocol explores the perceptual impact of ambiguity or incorrectness in the *output sequence*. The frameworks proposed for resolving prompt ambiguity, such as QA-TIED (Question-Answering Text-to-Image Editing and Disambiguation) or VS-TIED (Visual Setup Text-to-Image Editing and Disambiguation)⁸⁰, which involve user interaction to clarify intent, could inspire interactive mechanisms for users to *correct* or *clarify* perceived temporal errors or ambiguities within a generated sequence.

The "Meta-Perception Audit Loop" lens is central to this investigation. The core questions are: How do human observers perceive and react to these subtle temporal errors or ambiguities? Do they reliably notice them? Does the presence of such errors break their sense of realism, disrupt their understanding of the depicted process, or diminish their trust in the visualization? This constitutes an audit of their meta-perception—their perception of the perceptual quality and correctness of the generated visual narrative.

This type of testing is particularly crucial for **human-in-the-loop systems for scientific visualization and discovery**.⁵⁵ If users, especially domain experts, cannot reliably detect "slightly wrong" temporal states, or if such subtly incorrect states lead to misinterpretations of the underlying biology, then the generative model is unsuitable for high-stakes scientific applications. The successful use of human-in-the-loop generative models in fields like materials discovery¹⁰² underscores the importance of such validation in similarly critical domains.

Beyond general visual realism, for biological visuals, there exists a specific and critical dimension of "biological plausibility," especially within temporal sequences. Humans, particularly those with domain expertise (e.g., biologists, medical professionals), possess strong prior knowledge and expectations about how biological processes unfold over time. "Slightly wrong" renderings directly test the generative model's adherence to these domain-specific priors and, simultaneously, the human observer's ability to detect deviations from these expected patterns. This form of evaluation is more nuanced and domain-specific than simply asking, "Does this individual image look real?" It probes the model's understanding of temporal logic and biological causality.

6. Synthesizing the Framework: Insights from Prompt Archetypes and Strategic Upgrade Modules

The preceding analyses of individual prompts lay the groundwork for a synthesized understanding of the proposed research framework. This synthesis is achieved by examining the interplay of the defined "Prompt Archetypes" and by considering the strategic implications of the "Prompt Upgrade Module."

6.1. The Interplay of Prompt Archetypes

The "PROMPT ARCHETYPE MATRIX" provides a set of conceptual lenses that recur across the individual research prompts, creating a cohesive structure for the overall investigation. Their interplay is crucial for understanding the framework's depth and breadth.

- **Typological Core + Failure Stack (Applied in Prompts 1, 2):** This archetype, focused on mapping systems via their breakdown modes, is central to Prompt 1's critical dissection of existing perceptual metrics like LPIPS and DISTS. By identifying and analyzing their failure modes—such as LPIPS's shift sensitivity² and DISTS's potential to introduce artifacts during optimization⁶—the limitations of current approaches are laid bare. This archetype also informs the conceptualization of the "Organic Realism Prior" in Prompt 2. Defining what constitutes a "failure" to achieve organic realism (e.g., unnatural repetition, incorrect microstructures, lack of biological plausibility, similar to GAN failure modes like mode collapse¹⁰⁴ which are relevant to the proposed diffusion-GAN fusion) is essential for shaping the constraints and objectives of the prior itself. The previously discussed epistemic drift in metric development and the notion of a "perceptual uncanny valley" arising from metric overfitting are direct outcomes of applying this failure-focused analytical lens.
- **Sensory-Emotional Scaffold (Applied in Prompt 3):** This archetype underpins the core ambition of Prompt 3: to elevate "Material Prompting" beyond mere specification to a tool for achieving affective visual synthesis. The semiotic encoding of materials is designed to evoke not only visual and implied tactile sensations but also emotional responses. The study of semiotics in intersemiotic translation, particularly concerning plastic categories (form, color, texture) and modes of visual enunciation, provides a theoretical basis for understanding how material prompts can be crafted to achieve specific sensory and emotional impacts.⁶² Psychophysical studies on material appearance, which link specific image features (like highlight sharpness and contrast) to perceived material properties⁹¹, offer empirical grounding for how these sensory cues can be targeted through prompting. The development of a standardized syntax for material prompting and the exploration of metaphor versus exemplar use are both informed by the goal of building this sensory-emotional scaffold.
- **Prompt Infrastructure Design (Applied in Prompts 2, 3):** The development of a robust and flexible prompt infrastructure, including well-defined prompt ontologies, is crucial for realizing the advanced prompting strategies envisioned in Prompts 2 and 3. For Prompt 3, such an infrastructure is essential for implementing a standardized syntax for material prompting (e.g., `[]`). For Prompt 2, it is foundational for the "meta-control prompts" designed for adaptive material-level guidance. General prompt engineering principles, such as structuring prompts and ensuring clarity⁷⁶, and the structured parameter inputs seen in PBR material generation systems⁷⁰, highlight the need for and potential forms of such an infrastructure. The ability of meta-control prompts to facilitate interactive evolution of textures, for example, is entirely dependent on a well-designed underlying prompt system.
- **Semiotic Drift Simulator (Applied in Prompts 1, 4):** This lens aids in understanding how the meaning or interpretation of AI-related constructs can shift or degrade. In the context of Prompt 1, it helps analyze how the "meaning" of a perceptual metric's score can drift away from representing true perceptual quality, especially as models are optimized against potentially flawed metrics. For Prompt 4, this lens is relevant to understanding how the perceived meaning and coherence of a visual narrative (e.g., a temporal sequence of biological images) can drift or break down if temporal continuity is violated or if significant ambiguity is injected into the sequence. The choices made by AI in intersemiotic translation can inherently alter meaning⁶², and ambiguity in inputs or outputs can lead to significant shifts in interpretation.⁸⁰
- **Cognitive Friction Auditor (Applied in Prompts 2, 3):** This archetype emphasizes the importance of usability in advanced prompting systems. If achieving nuanced control over the "Organic Realism Prior" (Prompt 2) or mastering the semiotics of material prompting (Prompt 3) requires users to learn and deploy overly convoluted, unintuitive, or error-prone prompting strategies, it creates significant cognitive friction. This friction can hinder adoption and limit the practical utility of the generative system, regardless of its underlying power. Effective prompt engineering aims for clarity and ease of use⁷⁷, and ambiguity, if not managed or intentionally leveraged, can increase cognitive load for users trying to achieve their intended results.⁸⁰
- **Recursive Temporal Logic Synthesizer (Applied in Prompt 4):** This archetype directly embodies the core

technical challenge of Prompt 4: designing and implementing simulation protocols (such as the Temporal Graph + Narrative-of-Thought strategy) that can test and enforce logical temporal continuity in generated visual narratives. This involves concepts of recursive processes in AI ⁹⁷, the brain's use of generative models for prediction (as in Active Inference ⁹⁸), and awareness of potential failure modes like model collapse when systems become overly self-referential without sufficient grounding.⁹⁹ The investigations into dynamically conditioned realism, NoT as cognitive scaffolding, and sensitivity to biological plausibility all fall under this archetype.

Prompt Archetype	Purpose (from Matrix)	Primary Domain Mapping (from Matrix)	Key Supporting Research & Implications
Typological Core + Failure Stack	Map systems via breakdown modes	Deep prompting architecture	² (Metric failures); ¹⁰⁴ (GAN failures). <i>Implication:</i> Understanding failure modes (e.g., shift sensitivity, artifact generation) is crucial for designing robust metrics and priors like the ORP.
Sensory-Emotional Scaffold	Encode cross-modal affective cues	Material + perceptual prompting	⁶² (Semiotics of affect, psychophysics of material cues). <i>Implication:</i> Material prompts can be designed to target specific sensory (visual, tactile) and emotional responses by encoding relevant semiotic and perceptual cues.
Prompt Infrastructure Design	Modularize interface evolution + maintenance	Prompt ontology + governance	⁷⁰ (Prompt engineering principles, PBR system inputs). <i>Implication:</i> Standardized syntax and ontologies are vital for complex, adaptive prompting systems like meta-control prompts and nuanced material specification.
Semiotic Drift Simulator	Track meaning mutation through time or context	Psycholinguistic prompt feedback	⁶² (Meaning alteration in AI translation, impact of ambiguity). <i>Implication:</i> The interpretation of metric scores and visual narratives can degrade; prompts must manage ambiguity to maintain intended meaning.
Cognitive Friction Auditor	Surface prompt usability load for non-intuitive outputs	UX + prompt transparency	⁷⁷ (Prompt clarity, cognitive load from ambiguity). <i>Implication:</i> Advanced control mechanisms must be intuitive;

			high cognitive friction in prompting will limit the practical utility of sophisticated generative features.
Recursive Temporal Logic Synthesizer	Simulate state continuity over time	Time-aware visual narrative prompts	⁹⁸ (Active inference, risks of recursive collapse). <i>Implication:</i> Validating temporal coherence in generated sequences requires explicit modeling of temporal logic and awareness of self-referential system dynamics.

Table 2: Mapping Prompt Archetypes to Research Themes and Snippets

6.2. Leveraging Meta-Cycle Prompts for Iterative Advancement

The "PROMPT UPGRADE MODULE" provides meta-level prompts designed to challenge core assumptions and spur iterative refinement of the entire research framework.

Meta-Refinement Prompt: "Audit the assumptions baked into the notion of ‘Organic Realism’ across disciplines (medical illustration, perceptual psychophysics, cinematic VFX). What is elided or over-indexed in each domain’s definition of visual fidelity?"

This prompt forces a critical examination of what "realism" truly means in different contexts, which is vital for developing a nuanced "Organic Realism Prior" (ORP).

- Medical Illustration & Imaging Realism:** The primary goal here is often diagnostic accuracy and clarity for educational or clinical purposes. "Realism" is intrinsically linked to clinical utility and faithfulness to biological structure.¹⁰⁶ Quantitative evaluation might involve assessing the "clinical realism" and the similarity of statistical distributions between real and synthetic medical images.¹⁰⁶ Commercial systems like Konica Minolta's REALISM for X-rays emphasize depicting the entire image clearly while maintaining contrast, optimizing sharpness for specific sensor resolutions, and controlling granularity.¹⁰⁷
 - Aspects Elided:* Often, complex surface micro-details or stochastic variations are simplified if they are not diagnostically relevant. Full physical simulation of light transport is typically not a primary concern if clarity is achieved.
 - Aspects Over-indexed:* Clarity of key anatomical structures, contrast for differentiating tissues or pathologies, adherence to established anatomical representations and conventions.
- Perceptual Psychophysics of Realism:** This field defines realism as achieving visual indistinguishability from real-world scenes or objects.⁸⁷ The focus is on fulfilling essential visual requirements based on the known limitations and mechanisms of the human visual system (HVS). For instance, humans resolve visual ambiguities by making strong prior assumptions about scenes and objects.⁸⁹ Studies have shown that the perception of material properties in paintings and physically accurate renderings can be very similar if key visual cues (like highlight sharpness and contrast) are faithfully reproduced, even if other physical details differ.⁹²
 - Aspects Elided:* Absolute physical accuracy is not paramount if a perceptual match can be achieved through other means. Sometimes, higher-order semantic consistency might be less emphasized if the focus is on low-level visual cues.

- *Aspects Over-indexed*: Matching human perceptual thresholds, exploiting known properties of the HVS (e.g., binocular fusion for enhanced contrast in stereo displays⁸⁸), and ensuring the presence of critical visual cues necessary for material and shape perception.
- **Cinematic VFX Realism**: In visual effects, the aim is typically photorealism that serves a narrative purpose and enhances audience immersion. "Perceptual realism" in this context means creating visual experiences that align with our real-world perception sufficiently to make fictitious events, characters, or creatures believable within the story's world.¹⁰⁹ This can involve "invisible effects" that subtly enhance perceived reality or the creation of fantastical elements that nevertheless appear grounded in plausible physics or biology.
 - *Aspects Elided*: True physical correctness is often sacrificed if it does not look "right" on screen or if it detracts from the story or visual impact. VFX artists frequently "cheat" physics for aesthetic or narrative effect.
 - *Aspects Over-indexed*: Plausibility within the established diegesis of the story world, visual consistency across shots and sequences, overall "wow factor" and aesthetic appeal, and adherence to cinematic conventions of lighting, composition, and motion. The indexicality (trace of the real) and iconicity (resemblance to reality) of VFX elements are key considerations for their perceived truthfulness.¹⁰⁹

The concept of an **ontology of realism** itself is complex. One can distinguish between ontology as "saying what is" (a linguistic or conceptual representation) and ontography as "writing or recording what is" (a more direct, perhaps machine-mediated, inscription of reality).¹¹¹ Different disciplines effectively operate with different "ontologies" of realism. For example, in the context of creating realistic synthetic imagery of industrial facilities, realism extends beyond mere photo-realism to include the accurate representation of process flows and the functional characteristics of objects within the facility.¹¹²

"Organic Realism," therefore, is not a monolithic concept. Its specific "truth conditions" and the aspects of fidelity that are prioritized vary significantly across disciplines. Medical imaging values diagnostic truth and clarity above all. Psychophysics seeks perceptual indistinguishability, leveraging HVS characteristics. Cinematic VFX prioritizes narrative plausibility and visual impact. A successful and versatile ORP, as envisioned in Prompt 2, must either be flexible enough to accommodate these diverse "flavors" of realism—perhaps through tunable parameters or conditional generation—or the system must support multiple, domain-specific ORPs tailored to particular application requirements.

Discipline	Primary Goal of Realism	Key Metrics/Heuristics Used	Aspects Elided (Often Ignored/Simplified)	Aspects Over-Indexed (Highly Emphasized)
Medical Illustration/Imaging	Diagnostic accuracy, Clarity for education/treatment	Anatomical correctness, Ideal observer AUC ¹⁰⁶ , Clarity/Contrast ¹⁰⁷ , Pathological feature visibility	Full light physics if not diagnostic, Stochasticity if non-key, Complex micro-details if not informative	Key diagnostic features, Anatomical landmarks, Contrast between tissues, Adherence to medical conventions
Perceptual Psychophysics	Perceptual Indistinguishability from real world	Visual Turing Test ¹⁰⁸ , Just Noticeable Differences (JNDs), Match-to-sample tasks, Subjective rating scales	Absolute physical accuracy if perceptually matched, Some higher-order semantics if focused on	Critical visual cues for shape/material, HVS limitations & mechanisms, Perceptual thresholds,

		for realism/quality ⁹¹	low-level cues	Consistency
Cinematic VFX	Narrative Plausibility, Immersion, Visual Impact	Visual consistency, "Does it look right?", Adherence to physics (often selectively), Audience response, Believability within story context	Strict physical correctness if narratively unhelpful or visually unappealing, Invisible details	Visual impact, Narrative support, Emotional resonance, Cinematic aesthetics, Perceived weight and motion

Table 3: Cross-Disciplinary Perspectives on ‘Organic Realism’

Failure-Informed Inversion Prompt: "Design a generative pipeline that embraces the ‘Authenticity Gap’ instead of trying to close it. How can the ‘slight wrongness’ be recontextualized for educational clarity or affective resonance?"

This prompt challenges the conventional goal of perfect mimesis by suggesting that the "authenticity gap"—the perceptible differences between AI-generated content and human-created or real-world analogues—can be a source of value. AI-generated art and images often exhibit a kind of "hyperrealism" or an "unnatural precision" that can make them appear historically inconsistent or simply "off" when compared to traditional human artworks.¹¹³ There exists a persistent, albeit narrowing, gap between AI-generated and human-crafted work. For instance, AI's historical struggle with integrating legible text naturally within images has been cited as a failing in authenticity for applications like integrated branding.¹¹⁴

Instead of solely striving to eliminate this gap, one can "embrace the slight wrongness":

- **Educational Clarity:** Intentional abstraction, simplification, or even exaggeration of certain features—which constitute a form of "wrongness" from a purely photorealistic perspective—can significantly enhance pedagogical value. By strategically deviating from strict realism, key concepts or structures can be highlighted, making complex information more accessible and understandable (this connects directly to Prompt 2's exploration of balancing realism and pedagogical abstraction). Think of stylized medical diagrams that deliberately omit certain details to emphasize others.
- **Affective Resonance:** Many art forms, such as surrealism, expressionism, or impressionism, deliberately leverage "wrongness" (i.e., deviations from photographic reality) to evoke specific emotions, create a particular mood, or convey abstract ideas. The inherent "slight wrongness" in AI-generated visuals, if controlled and directed, could be steered towards these artistic ends, creating unique affective experiences.
- **Creative Exploration and Novel Aesthetics:** The "quirks," "hallucinations," and unexpected artifacts produced by generative models⁴⁷ need not always be viewed as errors. If deliberately controlled, refined, and contextualized, these deviations can become starting points for entirely novel aesthetics and creative styles that are uniquely characteristic of AI-assisted art.

Designing a generative pipeline to embrace this authenticity gap might involve:

- **Guided or Controllable Hallucination:** Training models not just to avoid hallucinations, but to generate specific types of "desired" artifacts or stylistic deviations in a controllable manner.
- **Difference Modeling or Stylistic Transformation Learning:** Training a model on pairs of (realistic_image, artistically_stylized_image) to explicitly learn the "transformation" function that introduces affectively resonant or pedagogically useful "wrongness."
- **Latent Space Navigation for Non-Realism:** Identifying and mapping regions in the model's latent space that correspond to these "interestingly wrong" or abstract aesthetics, and then developing methods to guide the

generation process towards or within these regions.

The concept of **constructive ambiguity**⁸⁰, while often discussed in the context of resolving ambiguity in input prompts, can be re-purposed here. The *output* of a generative model could possess "constructive ambiguity"—an intentional lack of specificity, a dream-like quality, or an invitation to multiple interpretations—that enhances affective resonance or stimulates deeper engagement.

What is currently perceived by some as a failure of authenticity¹¹³ could be reframed and cultivated as a unique stylistic signature. The "authenticity gap" is not merely a deficiency to be overcome but potentially a rich space for new aesthetic possibilities that are distinctly "AI-generated" rather than purely imitative of human art forms. This reframes a technical limitation or characteristic as a potential creative affordance, opening doors to new forms of visual expression.

Ontology Rewriter Prompt: "Rewrite your actor-agent ontology from the standpoint of semantic failure as first-class input. How would your pipeline change if ambiguity and hallucination were constructive layers instead of bugs?"

The current, often implicit, actor-agent ontology in generative AI typically assumes that actors (users) provide clear, unambiguous prompts, and agents (models) are expected to execute these prompts faithfully and accurately. Semantic failures, such as misinterpreting ambiguous prompts or generating hallucinatory content, are generally treated as bugs to be fixed or minimized.

Rewriting this ontology to treat semantic failure as a first-class input implies a fundamental shift:

- **Actors (Users):** Are empowered to intentionally provide ambiguous, vague, underspecified, or even contradictory prompts. The purpose is not to "trick" the model, but to use these inputs as probes to explore creative spaces, to elicit diverse interpretations, or to tap into the model's "imaginative" capabilities.
- **Agents (Models):** Must be equipped with new mechanisms beyond simple execution:
 - *Detect and Characterize Ambiguity/Hallucination:* The model should not just flag these instances but attempt to understand their type, source, and potential for creative exploration. While some work focuses on resolving ambiguity through clarification dialogues⁸⁰, a rewritten ontology might prioritize *leveraging* this ambiguity.
 - *Generate Diverse Interpretations:* Instead of striving for a single "best guess" interpretation of an ambiguous prompt, the model could offer multiple, distinct outputs that reflect different plausible resolutions of the ambiguity (this is akin to the Visual Setup T2I Editing and Disambiguation (VS-TIED) approach⁸⁰, but framed for creative exploration rather than disambiguation to a single correct answer).
 - *Controlled and Steerable Hallucination:* The model should allow users to "dial up" or "steer" hallucinatory tendencies along specific semantic or stylistic paths, transforming random artifacts into intentional creative elements.

This ontological shift would necessitate significant **pipeline changes**:

- **Prompt Interface:** Would need to evolve beyond simple text boxes. It might include tools for users to specify their tolerance for ambiguity, to choose preferred modes of resolving or exploring ambiguity (e.g., "surprise me with diverse interpretations," "show three distinct visual setups," "explore metaphorical connections"), or to provide feedback on which types of "failures" are interesting.
- **Model Architecture:** May require models specifically trained to generate diverse and creative outputs from underspecified or ambiguous inputs. Architectures that can explicitly represent, manipulate, and even seek out productive forms of uncertainty (related to epistemic and aleatoric uncertainty concepts³³) would be

beneficial.

- **Feedback Loops:** Iterative human-in-the-loop feedback⁵⁵ on which ambiguous interpretations or controlled hallucinations are deemed "interesting," "useful," or "aesthetically pleasing" could be used to refine the model's creative exploration capabilities and its understanding of productive semantic failures.

In this rewritten ontology, ambiguity and hallucination become **constructive layers**:

- **Ambiguity as Exploration:** An ambiguous prompt like "a creature of the deep forest, embodying both fear and wonder" becomes an explicit invitation for the model to generate a diverse portfolio of visual options, exploring different balances of these abstract concepts.
- **Hallucination as Novelty:** Controlled or guided hallucinations (e.g., "a cityscape where buildings are made of flowing water and light") can be a powerful source of novel artistic concepts and visual metaphors. As noted, hallucinations are not always failures and can lead to unexpected discoveries.⁴⁷

Reframing semantic failure in this way transforms the user-AI relationship from a master-servant dynamic to that of **dialogic partners in a creative process**. The AI doesn't just execute precise instructions; it explores, proposes, and even "misunderstands" in ways that can be creatively generative. The user's role shifts from being a pure specifier to becoming a curator, a guide, and a collaborator in this exploration of the AI's generative potential. This aligns with a vision of AI not just as a tool for producing predefined outputs, but as an instrument for augmenting human creativity and discovering new forms of expression.

7. Strategic Recommendations and Future Research Trajectories

The analyses conducted throughout this report, guided by the "Deep Research PROMPT SET," illuminate several critical areas where focused research and methodological shifts can significantly advance the field of generative AI. The following recommendations outline prioritized research questions, methodological considerations, and the potential impact of these endeavors.

7.1. Prioritized Research Questions

1. **Unified Perceptual Metrics:** How can a truly unified perceptual metric (conceptually aligned with MTAS/UniSim¹⁶) be developed that robustly correlates with human judgment across diverse generative tasks (e.g., image synthesis, texture generation, material appearance) and various model architectures? Such a metric must also demonstrate sensitivity to temporal dynamics (as explored in Prompt 4) and be inherently resistant to the kind of overfitting that plagues current metrics.³⁵
2. **Operationalizing the Organic Realism Prior (ORP):** What are the optimal architectural designs and training strategies for effectively operationalizing an "Organic Realism Prior" for biological textures? How can this prior be structured to allow a controllable balance between high-fidelity realism and pedagogical abstraction (Prompt 2)? Furthermore, how can meta-control prompts be designed to intuitively and effectively navigate and modulate this ORP?
3. **Standardized Material Prompting Syntax:** What constitutes the most effective, extensible, and user-friendly standardized syntax for "Material Prompting" (Prompt 3)? This syntax must balance rich semiotic expressiveness with computational parsability for generative models. A key aspect will be its ability to seamlessly integrate and reconcile both metaphorical and exemplar-based guidance for material specification.⁷⁸
4. **Scalable Psychophysical Validation of Temporal Coherence:** How can recursive perception simulations, incorporating Narrative-of-Thought strategies (Prompt 4), be scaled, standardized, and validated to provide reliable and efficient human-in-the-loop evaluation of temporal consistency and biological plausibility in complex generated visual sequences? This includes developing robust methodologies for ambiguity injection

and assessing its perceptual impact.

5. **Constructive Use of Semantic Failure:** What new generative model architectures, training paradigms, and interactive interfaces are required to constructively leverage semantic ambiguity and controlled hallucination as first-class inputs for creative exploration and affective resonance (as per the Ontology Rewriter Prompt)? This involves moving beyond treating these phenomena solely as bugs to be fixed.⁴⁷

7.2. Methodological Considerations

- **Expansion of Large-Scale Human Evaluation:** There is a clear need to expand the scope and scale of psychophysical studies⁷ to cover a broader range of generative stimuli (beyond simple images), more diverse perceptual tasks (including those related to material properties and temporal dynamics), and more complex temporal aspects of visual narratives. Standardized protocols for Human-in-the-Loop (HIL) evaluation, particularly in specialized scientific domains like biology and medicine, must be developed and disseminated.⁵⁵
- **Fostering Cross-Disciplinary Collaboration:** Addressing the complex challenges outlined requires deep collaboration between AI researchers and experts from diverse fields, including perceptual psychophysics, cognitive neuroscience, semiotics, linguistics, art and design, and specific scientific domains (e.g., biology, materials science, medicine). This collaboration should be integral to the design, development, and evaluation phases of new models and metrics.
- **Development of Advanced Benchmark Datasets:** The creation of new, challenging benchmark datasets is essential. These datasets should be specifically designed to test the nuanced aspects of generation and perception highlighted in the prompts, such as the temporal realism of biological processes, the affective resonance of synthesized materials, the ability to follow complex narrative prompts, and the capacity for constructive use of ambiguity. The UniSim-Bench¹⁶ provides a valuable starting point for perceptual similarity but needs to be augmented for these more advanced capabilities.
- **Emphasis on Interpretability and Disentanglement:** Future research should prioritize the development of generative models with more interpretable latent spaces and improved disentanglement of underlying perceptual factors.⁴⁹ Enhanced interpretability is key to understanding model behavior, diagnosing failures, and enabling more precise control over the generative process.

7.3. Potential Impact

Successfully addressing these research questions and adopting these methodological considerations holds the potential for transformative impact:

- **Paradigm Shift in Evaluation:** A move away from simplistic score-chasing towards a more holistic and meaningful assessment of generative AI, focused on genuine perceptual alignment, task-specific utility, and robustness.
- **Enhanced Generative Capabilities:** The development of more controllable, reliable, and versatile generative models capable of producing outputs that exhibit high fidelity, contextual awareness, and profound affective resonance. This will be particularly impactful in specialized domains such as scientific research, medical imaging, education, and artistic creation.
- **Novel Creative Tools and Expressions:** By learning to constructively leverage AI's inherent "quirks," ambiguities, and even "failures," new creative tools and paradigms can emerge, fostering novel forms of artistic expression and human-AI co-creation.
- **Increased Trust and Adoption in Critical Applications:** More robust validation methodologies, coupled with a deeper understanding of model capabilities, limitations, and failure modes, will be crucial for building trust and facilitating the responsible adoption of generative AI in critical applications where accuracy, reliability, and

ethical considerations are paramount.

In conclusion, the framework and inquiries presented within the "Deep Research PROMPT SET" offer a sophisticated roadmap for navigating the complex future of generative AI. By embracing critical evaluation, interdisciplinary approaches, and a willingness to challenge existing assumptions, the research community can steer the development of AI systems that are not only more powerful but also more aligned with human values, perception, and diverse needs.

Works cited

1. richzhang/PerceptualSimilarity: LPIPS metric. pip install lpips - GitHub, accessed June 10, 2025, <https://github.com/richzhang/PerceptualSimilarity>
2. Shift-tolerant Perceptual Similarity Metric - European Computer ..., accessed June 10, 2025, https://www.ecva.net/papers/eccv_2022/papers_ECCV/papers/136780089.pdf
3. (PDF) Shift-Tolerant Perceptual Similarity Metric - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/365037242_Shift-Tolerant_Perceptual_Similarity_Metric
4. A Review of the Image Quality Metrics used in Image Generative Models - Paperspace Blog, accessed June 10, 2025, <https://blog.paperspace.com/review-metrics-image-synthesis-models/>
5. Five Pitfalls When Assessing Synthetic Medical Images with Reference Metrics - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2408.06075v2>
6. Optimization results by the DISTS, LPIPS, and DeepWSD. The... | Download Scientific Diagram - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/figure/Optimization-results-by-the-DISTS-LPIPS-and-DeepWSD-The-noise-to-land-and-the_fig5_362568022
7. Exposing flaws of generative model evaluation metrics and their ..., accessed June 10, 2025, <https://openreview.net/forum?id=08zf7kTOoh-eld=SGRHI9AGow>
8. Exposing flaws of generative model evaluation metrics and their unfair treatment of diffusion models, accessed June 10, 2025, https://proceedings.neurips.cc/paper_files/paper/2023/hash/0bc795afae289ed465a65a3b4b1f4eb7-Abstract-Conference.html
9. Image Quality Assessment: Unifying Structure and Texture Similarity. - Matilda, accessed June 10, 2025, <https://matilda.science/work/e027b103-b85f-4b46-a6d8-66a08a949ad1>
10. Image Quality Assessment: Unifying Structure and Texture Similarity - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/347843713_Image_Quality_Assessment_Unifying_Structure_and_Texture_Similarity
11. (PDF) Image Quality Assessment: Unifying Structure and Texture Similarity - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/340683492_Image_Quality_Assessment_Unifying_Structure_and_Texture_Similarity
12. Image Quality Assessment: Unifying Structure and Texture Similarity - CityU

- Scholars, accessed June 10, 2025, [https://scholars.cityu.edu.hk/en/publications/image-quality-assessment\(aa061010-a46d-464a-91ef-8bf7fccb7080\).html](https://scholars.cityu.edu.hk/en/publications/image-quality-assessment(aa061010-a46d-464a-91ef-8bf7fccb7080).html)
13. dingkeyan93/DISTS: IQA: Deep Image Structure and ... - GitHub, accessed June 10, 2025, <https://github.com/dingkeyan93/DISTS>
 14. Adaptive Structure and Texture Similarity Metric for Image Quality Assessment and Optimization - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/375790946_Adaptive_Structure_and_Texture_Similarity_Metric_for_Image_Quality_Assessment_and_Optimization
 15. [Literature Review] MTA: Multimodal Task Alignment for BEV Perception and Captioning, accessed June 10, 2025, <https://www.themoonlight.io/en/review/mta-multimodal-task-alignment-for-bev-perception-and-captioning>
 16. Towards Unified Benchmark and Models for Multi-Modal Perceptual Metrics - OpenReview, accessed June 10, 2025, <https://openreview.net/pdf/d178a15aade2ad1061d54f18e281d284582dc55b.pdf>
 17. Towards Unified Benchmark and Models for Multi-Modal Perceptual Metrics - OpenReview, accessed June 10, 2025, <https://openreview.net/forum?id=EMAb263NIh>
 18. Towards Unified Benchmark and Models for Multi-Modal Perceptual Metrics - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2412.10594v1>
 19. Structural similarity index measure - Wikipedia, accessed June 10, 2025, https://en.wikipedia.org/wiki/Structural_similarity_index_measure
 20. Foundation Models Boost Low-Level Perceptual Similarity Metrics - PDXScholar, accessed June 10, 2025, https://pdxscholar.library.pdx.edu/cgi/viewcontent.cgi?article=1382&context=com_psci_fac
 21. Similarity and quality metrics for MR image-to-image translation - PMC, accessed June 10, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11785996/>
 22. Daily Papers - Hugging Face, accessed June 10, 2025, [https://huggingface.co/papers?q=Latent%20Consistency%20Model%20\(LCM\)](https://huggingface.co/papers?q=Latent%20Consistency%20Model%20(LCM))
 23. [2506.00327] Latent Guidance in Diffusion Models for Perceptual Evaluations - arXiv, accessed June 10, 2025, <https://arxiv.org/abs/2506.00327>
 24. (PDF) Latent Guidance in Diffusion Models for Perceptual Evaluations - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/392334412_Latent_Guidance_in_Diffusion_Models_for_Perceptual_Evaluations
 25. LGDM: Latent Guidance in Diffusion Models for Perceptual Evaluations - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2506.00327v1>
 26. Shreshth Saini, accessed June 10, 2025, <https://shreshthsaini.github.io/>
 27. Computer Vision and Pattern Recognition - arXiv, accessed June 10, 2025, <https://www.arxiv.org/list/cs.CV/recent?skip=318&show=500>
 28. The Unreasonable Effectiveness of Deep Features as a Perceptual Metric - arXiv, accessed June 10, 2025, <https://arxiv.org/abs/1801.03924>
 29. A Causal Framework for Aligning Image Quality Metrics and Deep Neural Network

- Robustness - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2503.02797v1>
30. A Survey on Image Quality Assessment: Insights, Analysis, and Future Outlook - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2502.08540v1>
 31. On the Epistemic Significance of Perceptual Structure - Oxford Academic, accessed June 10, 2025, <https://academic.oup.com/pq/article/74/1/1/7048481>
 32. Epistemological Foundations of Objectivist and Interpretivist Research - eCommons, accessed June 10, 2025, <https://ecommons.udayton.edu/cgi/viewcontent.cgi?article=1054&context=books>
 33. Uncertainty-Aware Regularization for Image-to-Image Translation - CVF Open Access, accessed June 10, 2025, https://openaccess.thecvf.com/content/WACV2025/papers/Vats_Uncertainty-Aware-Regularization-for-Image-to-Image-Translation_WACV_2025_paper.pdf
 34. Estimating Epistemic and Aleatoric Uncertainty with a Single Model - NIPS papers, accessed June 10, 2025, https://proceedings.neurips.cc/paper_files/paper/2024/file/c693c3ff83259aebcd55a41ab19a5d84-Paper-Conference.pdf
 35. A review on the generative models and its performance metrics, accessed June 10, 2025, <http://www.csam.or.kr/journal/view.html?uid=2195&&vmd=Full>
 36. Feature Likelihood Divergence: Evaluating the Generalization of Generative Models Using Samples - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2302.04440v4>
 37. A Pragmatic Note on Evaluating Generative Models with Fréchet Inception Distance for Retinal Image Synthesis - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2502.17160v1>
 38. Evaluating Generative Models for Tabular Data: Novel Metrics and Benchmarking - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2504.20900v1>
 39. A Good Score Does not Lead to A Good Generative Model - arXiv, accessed June 10, 2025, <https://arxiv.org/pdf/2401.04856>
 40. Five Pitfalls When Assessing Synthetic Medical Images with Reference Metrics | Request PDF - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/384764650_Five_Pitfalls_When_Assessing_Synthetic_Medical_Images_with_Reference_Metrics
 41. Similarity and quality metrics for MR image-to-image translation - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/388565272_Similarity_and_quality_metrics_for_MR_image-to-image_translation
 42. HYPE: Human eYe Perceptual Evaluation of Generative Models - OpenReview, accessed June 10, 2025, <https://openreview.net/forum?id=0zFjZ5nODa>
 43. Human Evaluation of Models Built for Interpretability - AAAI, accessed June 10, 2025, <https://aaai.org/ojs/index.php/HCOMP/article/download/5280/5132>
 44. Define your evaluation metrics | Generative AI on Vertex AI - Google Cloud, accessed June 10, 2025, <https://cloud.google.com/vertex-ai/generative-ai/docs/models/determine-eval>
 45. How can we evaluate generative language models? | Fast Data Science, accessed June 10, 2025,

- <https://fastdatascience.com/generative-ai/how-can-we-evaluate-generative-language-models/>
46. What Is Latent Space? - IBM, accessed June 10, 2025, <https://www.ibm.com/think/topics/latent-space>
 47. The Secret Lives of Generative AI Models: What They Create When We're Not Watching, accessed June 10, 2025, <https://www.computer.org/publications/tech-news/trends/secret-lives-gen-ai>
 48. Realism Index: Interpolation in Generative Models With Arbitrary Prior | OpenReview, accessed June 10, 2025, <https://openreview.net/forum?id=SkIVqa4YwH>
 49. Disentangled Representation Learning - IEEE Computer Society, accessed June 10, 2025, <https://www.computer.org/csdl/journal/tp/2024/12/10579040/1YeeHU2QTkc>
 50. Disentangled Representation Learning - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2211.11695v4>
 51. Disentangling Generative Factors of Physical Fields Using Variational Autoencoders - Frontiers, accessed June 10, 2025, <https://www.frontiersin.org/journals/physics/articles/10.3389/fphy.2022.890910/full>
 52. Disentangled Representation Learning of Deep Generative Models | PPT - SlideShare, accessed June 10, 2025, <https://www.slideshare.net/quolc/disentangled-representation-learning-of-deep-generative-models>
 53. Material transforms from disentangled NeRF representations - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2411.08037v1>
 54. Explicit Disentanglement of Appearance and Perspective in Generative Models - NIPS papers, accessed June 10, 2025, <http://papers.neurips.cc/paper/8387-explicit-disentanglement-of-appearance-and-perspective-in-generative-models.pdf>
 55. Importance of Human-in-the-Loop for Generative AI: Balancing Ethics and Innovation, accessed June 10, 2025, <https://www.digitaldividedata.com/blog/human-in-the-loop-for-generative-ai>
 56. Understanding Human Evaluation Metrics in AI: What They Are and How They Work, accessed June 10, 2025, <https://galileo.ai/blog/human-evaluation-metrics-ai>
 57. Discrete Latent Graph Generative Modeling with Diffusion Bridges - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2403.16883v1>
 58. [2409.07912] Multi-granularity Score-based Generative Framework Enables Efficient Inverse Design of Complex Organics - arXiv, accessed June 10, 2025, <https://arxiv.org/abs/2409.07912>
 59. Multiclass Texture Synthesis Using Generative Adversarial Networks - SciTePress, accessed June 10, 2025, <https://www.scitepress.org/Papers/2023/116575/116575.pdf>
 60. Generating Infinite-Size Textures using GANs with Patch-by-Patch Paradigm - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2309.02340v3>
 61. Image Quality Assessment: Exploring Joint Degradation Effect of Deep Network

- Features via Kernel Representation Similarity Analysis - IEEE Computer Society, accessed June 10, 2025, <https://www.computer.org/csdl/journal/tp/2025/04/10833846/23IjUuSUUUM>
62. For a Semiotic Approach to Generative Image AI: On ... - ORBi, accessed June 10, 2025, <https://orbi.uliege.be/bitstream/2268/331490/1/SR-9-Big-Dada-2025.pdf>
 63. Signata 17 / Multimodal Translation in Generative AI: Fidelity and Creativity, accessed June 10, 2025, <https://journals.openedition.org/signata/5396>
 64. Semiotics and Artificial Intelligence (AI): An Analysis of Symbolic Communication in the Age of Technology - ResearchGate, accessed June 10, 2025, https://www.researchgate.net/publication/379020156_Semiotics_and_Artificial_Intelligence_AI_An_Analysis_of_Symbolic_Communication_in_the_Age_of_Technology
 65. An inter-semiotic analysis of ideational meaning in text-prompted AI-generated images, accessed June 10, 2025, https://www.researchgate.net/publication/377489601_An_inter-semiotic_analysis_of_ideational_meaning_in_text-prompted_AI-generated_images
 66. Talking to Machines: Semiotic Analysis, Implications for Teaching and Media Literacy, accessed June 10, 2025, <https://riviste.unimi.it/index.php/anicon/article/view/23944>
 67. "Prompt Engineering" is literally just communication skills : r/aiwars - Reddit, accessed June 10, 2025, https://www.reddit.com/r/aiwars/comments/1jq6j9v/prompt_engineering_is_literally_just/
 68. [ICCV 2023] Text2Tex: Text-driven Texture Synthesis via Diffusion Models - GitHub, accessed June 10, 2025, <https://github.com/daveredrum/Text2Tex>
 69. Text2Tex: Text-driven Texture Synthesis via Diffusion Models - Dave Zhenyu Chen, accessed June 10, 2025, <https://daveredrum.github.io/Text2Tex/>
 70. giuvecchio/matfuse-sd: MatFuse: Controllable Material ... - GitHub, accessed June 10, 2025, <https://github.com/giuvecchio/matfuse-sd>
 71. MaterialFusion: High-Quality, Zero-Shot, and Controllable Material Transfer with Diffusion Models - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2502.06606v1>
 72. DreamMat: High-quality PBR Material Generation with Geometry ..., accessed June 10, 2025, <https://zzzyuqing.github.io/dreammat.github.io/>
 73. I made a seamless pbr material generator : r/StableDiffusion - Reddit, accessed June 10, 2025, https://www.reddit.com/r/StableDiffusion/comments/1ic304r/i_made_a_seamless_pbr_material_generator/
 74. MaterialMVP: Illumination-Invariant Material Generation via Multi-view PBR Diffusion - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2503.10289v1>
 75. Mastering PBR Materials | Babylon.js Documentation, accessed June 10, 2025, <https://doc.babylonjs.com/features/featuresDeepDive/materials/using/masterPBR>
 76. Prompting Generative Models - Yext Hitchhikers, accessed June 10, 2025, <https://hitchhikers.yext.com/guides/write-prompts-for-generative-models/01-prompt-generative-models/>

77. Prompt Engineering: The Art of Getting What You Need From Generative AI, accessed June 10, 2025, <https://iac.gatech.edu/featured-news/2024/02/AI-prompt-engineering-ChatGPT>
78. [2503.17987v1] Metaphor-based Jailbreaking Attacks on Text-to-Image Models - arXiv, accessed June 10, 2025, <http://www.arxiv.org/abs/2503.17987v1>
79. Multimodal Chain-of-Thought Prompting for Metaphor Generation - ACL Anthology, accessed June 10, 2025, <https://aclanthology.org/2024.clicit-1.62.pdf>
80. Resolving Ambiguities in Text-to-Image Generative Models - Amazon Science, accessed June 10, 2025, <https://assets.amazon.science/91/b9/affc407f4fda8ce0b4d7430e9c1c/resolving-ambiguities-in-text-to-image-generative-models.pdf>
81. Grounding overview | Generative AI on Vertex AI - Google Cloud, accessed June 10, 2025, <https://cloud.google.com/vertex-ai/generative-ai/docs/grounding/overview>
82. Prompt engineering and the responsible use of generative AI - Fivetran, accessed June 10, 2025, <https://www.fivetran.com/blog/prompt-engineering-and-the-responsible-use-of-generative-ai>
83. Semantic and Sentiment Dual-Enhanced Generative Model for Script Event Prediction, accessed June 10, 2025, <https://aclanthology.org/2025.coling-main.622/>
84. Interpreting the Latent Space of GANs for Semantic Face Editing - CVF Open Access, accessed June 10, 2025, https://openaccess.thecvf.com/content_CVPR_2020/papers/Shen_Interpreting_the_Latent_Space_of_GANs_for_Semantic_Face_Editing_CVPR_2020_paper.pdf
85. LatentEditor: Text Driven Local Editing of 3D Scenes, accessed June 10, 2025, <https://latenteditor.github.io/>
86. FocalDreamer: Text-Driven 3D Editing via Focal-Fusion Assembly, accessed June 10, 2025, <https://focaldreamer.github.io/>
87. Human Perception of Image Visual Realism - N-CRIPT, accessed June 10, 2025, <https://ncrypt.comp.nus.edu.sg/site/ncrypt-top/visualrealism/>
88. Path from Photorealism to Perceptual Realism - Apollo, accessed June 10, 2025, <https://www.repository.cam.ac.uk/items/b4116c6a-ebdc-44a9-8e6f-cbc746fe8d5e>
89. Measuring Visual Shape using Computer Graphics Psychophysics, accessed June 10, 2025, https://cim.mcgill.ca/~langer/MY_PAPERS/LangerBuelthoff-EGWR00.pdf
90. The Role of Perception for Computer Graphics - Eurographics, accessed June 10, 2025, <https://diglib.eg.org/server/api/core/bitstreams/097f365f-799f-4449-ab69-2782c7104b60/content>
91. Probing the Link Between Vision and Language in Material Perception Using Psychophysics and Unsupervised Learning, accessed June 10, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10849714/>
92. Perception of material appearance: A comparison between painted and rendered images, accessed June 10, 2025,

- <https://pmc.ncbi.nlm.nih.gov/articles/PMC8131993/>
93. A Survey: Spatiotemporal Consistency in Video Generation - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2502.17863v1>
 94. Temporal Consistency in image models: Is 'Scene Memory' Possible? : r/StableDiffusion, accessed June 10, 2025, https://www.reddit.com/r/StableDiffusion/comments/1krr61k/temporal_consistency_in_image_models_is_scene/
 95. Analyzing temporal coherence for deepfake video detection - AIMS Press, accessed June 10, 2025, <https://www.aimspress.com/article/doi/10.3934/era.2024119?viewType=HTML>
 96. A Physical Coherence Benchmark for Evaluating Video Generation Models via Optical Flow-guided Frame Prediction - arXiv, accessed June 10, 2025, <https://arxiv.org/html/2502.05503v2>
 97. Recursive Generative Emergence, accessed June 10, 2025, <https://rgemergence.com/>
 98. Generative Models for Active Vision - Frontiers, accessed June 10, 2025, <https://www.frontiersin.org/journals/neurobotics/articles/10.3389/fnbot.2021.651432/full>
 99. What happens when generative AI models train recursively on each others' generated outputs? - arXiv, accessed June 10, 2025, <https://arxiv.org/pdf/2505.21677>
 100. V Generative Recursion, accessed June 10, 2025, https://felleisen.org/matthias/HtDP2e/part_five.html
 101. OVERCOMING THE AMBIGUITY REQUIREMENT USING GENERATIVE AI - MDU, accessed June 10, 2025, <https://mdu.diva-portal.org/smash/get/diva2:1931301/FULLTEXT01.pdf>
 102. Ternary materials discovery using human-in-the-loop generative machine learning, accessed June 10, 2025, <https://pubs.rsc.org/en/content/articlehtml/2025/ra/d5ra00427f>
 103. Ternary materials discovery using human-in-the-loop generative machine learning, accessed June 10, 2025, <https://pubmed.ncbi.nlm.nih.gov/40476230/>
 104. GANs Failure Modes: How to Identify and Monitor Them, accessed June 10, 2025, <https://neptune.ai/blog/gan-failure-modes>
 105. GANs: Understanding the Latest Trends in Generative Adversarial Networks - Netguru, accessed June 10, 2025, <https://www.netguru.com/blog/generative-adversarial-networks>
 106. Observer-study-based approaches to quantitatively evaluate the realism of synthetic medical images - PubMed Central, accessed June 10, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10411234/>
 107. Realism-xray - Healthcare | KONICA MINOLTA, accessed June 10, 2025, <https://www.konicaminolta.com/global-en/healthcare/technology/xray/realism/index.html>
 108. Path from Photorealism to Perceptual Realism - Department of Computer Science and Technology |, accessed June 10, 2025, https://www.cl.cam.ac.uk/~fz261/thesis/Dissertation_Path_from_Photorealism_to

[Perceptual_Realism_compressed.pdf](#)

109. Current Trends of VFX | Digital Sketchbook - UWL CampusPress, accessed June 10, 2025, <https://campuspress.uwl.ac.uk/weronika/current-trends-of-vfx/>
110. Principles of VFX | Digital Sketchbook - UWL CampusPress, accessed June 10, 2025, <https://campuspress.uwl.ac.uk/noor/principles-of-vfx/>
111. Ontology and Ontography in Digital Imaging - media/rep, accessed June 10, 2025, <https://mediarep.org/bitstreams/0bd6967c-dc2b-48f1-aaa7-565130a96162/download>
112. Use of Ontology to Facilitate the Creation of Synthetic Imagery of Industrial Facilities - CEUR-WS.org, accessed June 10, 2025, https://ceur-ws.org/Vol-808/STIDS2011_CR_P3_PopeRoberts.pdf
113. A Critical Assessment of Modern Generative Models' Ability to Replicate Artistic Styles, accessed June 10, 2025, <https://arxiv.org/html/2502.15856v1>
114. The Future of AI-Driven Design: From Generative Art to Integrated Branding | JAXONLABS, accessed June 10, 2025, <https://www.jaxonlabs.com/the-future-of-ai-driven-design-from-generative-art-to-integrated-branding>