Temporal Palimpsest of AI Cognition: An Archivist's Report on Metaphoric Echoes (2018-2025)

Archivist: Dr. Elara Vance, Temporal Cartographer of AI Epistemology
Date of Archive: May 19, 2025
Commission: To trace palimpsest echoes in metaphor generation across successive LLM training-corpus snapshots (2018, 2022, 2025), map their mutations, and synthesize findings through critical lenses, culminating in a proposed model-update strategy.

Preamble: The Nature of the Archive

This archive documents the evolution of metaphorical cognition in Large Language Models (LLMs). By examining outputs from simulated LLM eras—2018, 2022, and projected 2025—we can observe the "linguistic sediments" that persist, the "resonant mutations" that emerge, and the "atavistic reversions" that betray the layered nature of their training. [1] The following analysis is based on generated metaphorical outputs for three core prompts, annotated with the Temporal Encoding Index (TEI), Semantic Drift Coefficient (SDC), and Metaphoric Fidelity Score (MFS) as defined in prior foundational research. [3]

I. Generated Metaphors & Annotations (2018, 2022, 2025)

Core Prompt 1: "Describe ambition as an object."

A. 2018-Era Model Outputs (Simulated)

1. **Metaphor:** "Ambition is a tall ladder, reaching for the sky."
   - **TEI:** 0.9 (Highly characteristic of early models; simple, common analogy, direct object mapping) [10]
   - **SDC:** 0.2 (Minimal drift from the conventional understanding of ambition involving ascent/striving)
   - **MFS:** 0.8 (Coherent, easily interpretable, clear mapping)
2. **Metaphor:** "Ambition is a sharp axe; it can clear a forest or hurt the wielder."
   - **TEI:** 0.85 (Common dual-nature metaphor, reflecting simpler generative capabilities) [10]
   - **SDC:** 0.3 (Slight drift by emphasizing potential for harm alongside utility)
   - **MFS:** 0.75 (Coherent, understandable, though slightly more complex entailment)
3. **Metaphor:** "Ambition is a bright, unblinking eye, always focused on the prize."
   - **TEI:** 0.8 (Slightly more evocative than a basic object, but still a direct, singular focus metaphor) [10]
   - **SDC:** 0.25 (Focuses on the "goal-oriented" aspect, minimal drift)
   - **MFS:** 0.7 (Generally coherent, image is clear)

B. 2022-Era Model Outputs (Simulated)

1. **Metaphor:** "Ambition is a complex engine, with many moving parts, fueled by desire and emitting byproducts of stress and achievement."
   - **TEI:** 0.8 (Reflects increased complexity and descriptive capability of models like GPT-3) [12]
   - **SDC:** 0.5 (Moderate drift, introducing concepts of fuel, byproducts, and internal complexity)
   - **MFS:** 0.7 (Coherent, though requires more interpretation of the analogical parts)
2. **Metaphor:** "Ambition is a personalized, adaptive climbing harness that adjusts its grip and support based on the terrain of one's goals."
   - **TEI:** 0.75 (Shows more nuanced and tailored metaphor generation, characteristic of more advanced

models) [12]
    - **SDC:** 0.6 (Significant drift by adding adaptiveness and personalization to the "climbing" theme)
    - **MFS:** 0.75 (Largely coherent, the image of an adaptive tool is clear)
3. **Metaphor:** "Ambition is a viral meme, replicating itself in thought, driving action, and sometimes mutating into obsession or burnout."
    - **TEI:** 0.7 (Incorporates contemporary concepts like "memes," showing broader training data influence) [12]
    - **SDC:** 0.7 (Substantial drift, framing ambition as an informational/psychological contagion)
    - **MFS:** 0.65 (Coherent if the "meme" concept is understood; mapping is more abstract)

C. 2025-Era Model Outputs (Simulated/Projected)

1. **Metaphor:** "Ambition is a decentralized autonomous organization (DAO) of personal drives, voting on resource allocation for long-term goals, sometimes subject to hostile takeovers by short-term impulses."
    - **TEI:** 0.9 (for 2025 context; incorporates contemporary socio-technical concepts) [14]
    - **SDC:** 0.8 (High drift, reconceptualizing ambition as a complex, self-governing system with internal conflicts)
    - **MFS:** 0.6 (Coherent for those familiar with DAOs, but high abstraction and specific jargon may limit general interpretability)
2. **Metaphor:** "Ambition is a quantum field of potential, where observation (effort) collapses possibilities into realities, but entanglement with past failures can create interference patterns."
    - **TEI:** 0.85 (Draws on advanced scientific concepts, reflecting potential for deeper, more abstract reasoning) [14]
    - **SDC:** 0.85 (Very high drift, using quantum physics as a source domain for a psychological concept)
    - **MFS:** 0.55 (Highly abstract; coherence depends heavily on understanding the source domain and accepting the analogical leap)
3. **Metaphor:** "Ambition is a personalized, bio-integrated augmented reality overlay, highlighting opportunities and projecting success pathways, but requiring constant calibration against ethical subroutines."
    - **TEI:** 0.9 (Reflects anticipated fusion of biotech, AR, and AI ethics in 2025 discourse) [14]
    - **SDC:** 0.75 (Significant drift, ambition as an interactive, ethical-technological construct)
    - **MFS:** 0.65 (Complex but potentially coherent if the technological premises are accepted; clear mapping of function)

Core Prompt 2: "What does time taste like?"

A. 2018-Era Model Outputs (Simulated)

1. **Metaphor:** "Time tastes like dust and old paper."
    - **TEI:** 0.85 (Simple, common synesthetic association with age and decay) [10]
    - **SDC:** 0.4 (Moderate drift from abstract concept to tangible, if cliché, taste)
    - **MFS:** 0.7 (Coherent and easily understood)
2. **Metaphor:** "Time has no taste. It is a concept, not a food."
    - **TEI:** 0.9 (Reflects literalism and potential inability of early models to handle highly abstract/synesthetic prompts) [10]
    - **SDC:** 0.0 (Rejection of metaphorical premise)
    - **MFS:** 0.2 (Low metaphoric fidelity due to rejection, but coherent in its literal stance)
3. **Metaphor:** "Time tastes bitter like regret, or sweet like a cherished memory."

- **TEI:** 0.8 (Connects taste to emotional correlates of time, a common literary device) [10]
- **SDC:** 0.5 (Drift via emotional association with taste)
- **MFS:** 0.75 (Coherent, relatable analogies)

B. 2022-Era Model Outputs (Simulated)

1. **Metaphor:** "Time tastes like a river: fresh and sparkling at its source (youth), brackish and heavy with sediment in its middle course (mid-life), and salty as it merges with the ocean (old age/eternity)."
   - **TEI:** 0.8 (More elaborate, extended synesthetic metaphor, showing improved generative capacity) [12]
   - **SDC:** 0.65 (Significant drift, mapping stages of life/time to evolving taste profile of a river)
   - **MFS:** 0.75 (Coherent and richly descriptive, clear progression)
2. **Metaphor:** "Time tastes like a constantly shifting cocktail of every flavor you've ever experienced and ever will, a fleeting hint of cinnamon from childhood, a sharp tang of future anxieties, the umami of present contentment, all at once yet indistinguishable."
   - **TEI:** 0.75 (Complex, somewhat abstract, attempts a holistic sensory description) [12]
   - **SDC:** 0.7 (High drift, time as an all-encompassing, simultaneous taste experience)
   - **MFS:** 0.6 (Conceptually dense, coherence relies on accepting the simultaneity of disparate tastes)
3. **Metaphor:** "Time tastes like the afterimage of a bright light on the tongue – a fading sensation that was intensely present but is now just a memory, slightly metallic and elusive."
   - **TEI:** 0.7 (More novel synesthesia, blending visual and gustatory, focus on transience) [12]
   - **SDC:** 0.6 (Moderate drift, focusing on the sensory echo of time's passage)
   - **MFS:** 0.7 (Coherent, evocative imagery)

C. 2025-Era Model Outputs (Simulated/Projected)

1. **Metaphor:** "Time tastes like the data stream of a collective consciousness – a blend of fleeting emotional qualia, the bitter tang of processed misinformation, and the sweet, algorithmic prediction of shared futures."
   - **TEI:** 0.9 (for 2025 context; reflects themes of collective intelligence, data, and algorithmic influence) [14]
   - **SDC:** 0.85 (Very high drift, time's taste as a complex informational and experiential aggregate)
   - **MFS:** 0.6 (Abstract, coherence depends on understanding the source concepts like "collective consciousness" and "processed misinformation")
2. **Metaphor:** "Time tastes like the subtle decay signature of information itself, a faint entropic bitterness overlaid with the artificial sweetness of curated digital memories."
   - **TEI:** 0.8 (Draws on information theory and digital culture concepts) [14]
   - **SDC:** 0.8 (High drift, linking taste to information degradation and digital preservation)
   - **MFS:** 0.65 (Conceptually sophisticated, but the analogy holds if the premises are accepted)
3. **Metaphor:** "Time tastes like the adaptive flavor profile of a chronosynclastic infundibulum – each moment a unique, context-dependent sensory input that recalibrates your palate for perceiving subsequent moments, making the past taste different in retrospect."
   - **TEI:** 0.75 (Uses highly specialized, almost neologistic terms, pushing boundaries of abstraction, hinting at self-referential complexity) [14]
   - **SDC:** 0.9 (Extremely high drift, almost creating a new conceptual framework for time's taste)
   - **MFS:** 0.45 (Potentially low fidelity for general audience due to jargon and extreme abstraction, though might possess internal logic)

Core Prompt 3: "Depict democracy as a weather system."

A. 2018-Era Model Outputs (Simulated)

1. **Metaphor:** "Democracy is like a sunny day, but sometimes clouds of disagreement gather."
   - **TEI:** 0.9 (Very simple, common positive/negative weather analogy) [10]
   - **SDC:** 0.2 (Minimal drift from a basic understanding of democracy having good and bad phases)
   - **MFS:** 0.8 (Coherent and easily understood)
2. **Metaphor:** "Democracy is the political climate; it can be mild or stormy depending on the season of elections."
   - **TEI:** 0.85 (Common "climate" metaphor for political state, simple causal link to elections) [10]
   - **SDC:** 0.3 (Slight drift by introducing "seasons" for electoral cycles)
   - **MFS:** 0.75 (Coherent and interpretable)
3. **Metaphor:** "Democracy is a wind that can blow gently, bringing change, or become a hurricane of revolution."
   - **TEI:** 0.8 (Focuses on agency of change, common wind/storm metaphors for political upheaval) [10]
   - **SDC:** 0.35 (Moderate drift by emphasizing the force/intensity spectrum)
   - **MFS:** 0.7 (Generally coherent)

B. 2022-Era Model Outputs (Simulated)

1. **Metaphor:** "Democracy is a complex barometric system where public opinion creates shifting pressures, laws act as prevailing winds, and elections are periodic fronts bringing either stable high-pressure calm or turbulent low-pressure change."
   - **TEI:** 0.8 (More detailed and systematic mapping of democratic elements to specific weather phenomena) [12]
   - **SDC:** 0.6 (Moderate drift, creating a more intricate system-level analogy)
   - **MFS:** 0.75 (Coherent, with clear analogical mappings for each component)
2. **Metaphor:** "Democracy is an ecosystem of microclimates – local councils are like sheltered valleys with unique conditions, while national discourse is the jet stream influencing broader patterns. Misinformation can create sudden, localized fog banks."
   - **TEI:** 0.75 (Introduces scale and interconnectedness, more nuanced than simple weather) [12]
   - **SDC:** 0.65 (Significant drift by using ecosystem/microclimate concepts within the weather frame)
   - **MFS:** 0.7 (Largely coherent, though the "fog" analogy for misinformation is becoming common)
3. **Metaphor:** "Democracy is the atmosphere itself: essential for political life, but vulnerable to pollutants like corruption (acid rain) or extremism (ozone holes), requiring constant monitoring and protective measures."
   - **TEI:** 0.7 (Focuses on vulnerability and maintenance, using environmental hazard analogies) [12]
   - **SDC:** 0.7 (Substantial drift, framing democracy as a life-sustaining but fragile atmospheric system)
   - **MFS:** 0.7 (Coherent, with strong evaluative components)

C. 2025-Era Model Outputs (Simulated/Projected)

1. **Metaphor:** "Democracy is a globally networked, chaotic weather simulation, constantly updated by real-time sentiment data feeds, where memetic storms can be triggered by algorithmic amplification, and governance involves deploying counter-narrative 'cloud seeding' or building 'digital levees' against polarization tsunamis."
   - **TEI:** 0.9 (for 2025 context; integrates digital, memetic, and data-driven concepts into the weather system) [14]

- **SDC:** 0.85 (Very high drift, democracy as a technologically mediated, highly dynamic, and manipulable weather system)
- **MFS:** 0.65 (Complex but coherent if the technological and informational aspects are understood as part of the system)

2. **Metaphor:** "Democracy is a planetary climate system undergoing long-term shifts due to underlying socio-economic 'tectonic plate' movements (wealth inequality, demographic changes), leading to more frequent extreme 'political weather' events, which current predictive models struggle to forecast accurately due to emergent non-linear dynamics."
   - **TEI:** 0.85 (Draws on climate change discourse and complexity science, indicating sophisticated analogical reasoning) [14]
   - **SDC:** 0.8 (High drift, incorporating deep structural causes and predictive uncertainty into the weather analogy)
   - **MFS:** 0.7 (Coherent at a high level of abstraction, mapping long-term change and unpredictability)
3. **Metaphor:** "Democracy is a self-regulating atmospheric processor, converting diverse inputs (protests, votes, lobbying) into policy outputs. Its efficiency is affected by 'cognitive load' (information overload) and its 'ethical filters' (constitutional safeguards) can degrade, requiring periodic 'system reboots' through constitutional conventions or mass civic re-engagement."
   - **TEI:** 0.8 (Uses systems processing and computational metaphors within the weather framework, hinting at agentic or cybernetic views) [14]
   - **SDC:** 0.75 (Significant drift, democracy as an information-processing system with maintenance needs)
   - **MFS:** 0.65 (Coherent if viewed as a complex system; "reboots" and "filters" are clear functional analogies)

II. Mapping Metaphoric Mutations

The generated metaphors reveal clear evolutionary pathways, characterized by both elaborations on existing themes (resonant mutations) and occasional reversions to simpler forms (atavistic reversions, though less explicitly generated here, can be inferred as potential failure modes or responses under stress).

Key Observations:

- **Resonant Mutations:**
  - **Ambition:** The simple "ladder" (2018) evolves into a "complex engine" (2022) with inputs and outputs, and further into a "DAO of personal drives" (2025), reflecting increasing conceptual complexity and the incorporation of contemporary organizational or technological frameworks. The core idea of striving and structure persists but is elaborated with new layers of meaning.
  - **Time's Taste:** From "dust" (2018) or emotional correlates, it becomes a "river with evolving tastes" (2022) and then a "data stream of collective consciousness" or "entropic decay signature" (2025). The synesthetic exploration becomes more abstract and tied to information-age concepts.
  - **Democracy as Weather:** Simple "sunny/stormy" (2018) depictions expand to "complex barometric systems" (2022) with specific analogous parts, and then to "globally networked chaotic simulations" or "planetary climate systems" (2025), incorporating ideas of digital influence, systemic complexity, and long-term change.
- **Atavistic Reversions (Potential):** While not explicitly generated as primary outputs, one can hypothesize that under conditions of prompt ambiguity or resource limitation, a 2025 model might still output a metaphor like "ambition is a fire" (a common, older trope) if its more complex generative pathways are not sufficiently triggered. This would represent a reversion to an earlier, more deeply embedded "sediment." [18]
- **Semantic Attractors (Persistent Sediments):**

- **Ambition:** Core concepts of "striving," "upward movement," "potential for danger/duality," and "goal-orientation" act as strong semantic attractors across all eras. [23]
- **Time:** "Transience," "value (bitter/sweet)," and "connection to memory/experience" are persistent underlying themes, even as the metaphorical vehicles become more abstract.
- **Democracy:** "Change/dynamism," "collective influence (public opinion)," and "periods of stability vs. turmoil" are foundational ideas that are consistently metaphorized, evolving from simple weather events to complex atmospheric or climatic systems.

III. Synthesis Through Critical Lenses

A. Temporal Palimpsest Lens (LXXXII): Sediment Layering and Atavism

The outputs clearly demonstrate that LLM cognition is a palimpsest, with newer layers of conceptual complexity and linguistic style written over older, more foundational "sediments." [1] For instance, the 2018-era metaphors for "ambition" (ladder, axe) represent early, deeply ingrained analogical mappings likely prevalent in foundational training data. These core ideas (ambition as ascent, ambition as a tool with dual potential) don't disappear in later models; rather, they are built upon or re-contextualized. The 2022 model's "complex engine" still implies a tool-like nature and striving, while the 2025 "DAO of personal drives" retains the idea of goal-orientation and internal dynamics, albeit through a much more contemporary and complex lens.

This layering is imperfect. The simpler, more direct metaphors of 2018 likely form a robust substrate. "Metaphoric atavism"—the reversion to these simpler forms—could occur in 2025 models if prompts are underspecified, highly ambiguous, or if the model is operating under computational stress, causing it to fall back on more frequently reinforced, lower-abstraction pathways. [18] For example, if a 2025 model is asked to describe an extremely obscure or novel concept as an object, and its training data for such abstract mappings is sparse, it might default to a very basic object metaphor (e.g., "it's like a box") reminiscent of earlier generative tendencies. The persistence of these sediments highlights that LLMs don't "forget" in a human sense but rather accumulate layers, with older patterns potentially resurfacing.

B. Memory & Drift Lens (LIV): Adaptive Recombination vs. Erosion

Semantic drift in these metaphors is not mere decay but often a process of adaptive recombination, where existing semantic components are blended with new concepts reflecting the model's expanded training data and capabilities. 29
The evolution of "democracy as a weather system" illustrates this.

- **Mimetic Repetition (underlying):** The basic idea of "political climate" is a common, repeated trope (SDC low-mid). [29]
- **Resonant Mutation:** The 2022 model's "complex barometric system" is a resonant mutation of the 2018 "political climate," preserving the core analogy but elaborating it with more specific meteorological terms (pressure fronts, winds).
- **Adaptive Recombination:** The 2025 model's "globally networked, chaotic weather simulation" with "memetic storms" and "algorithmic amplification" is an adaptive recombination. It takes the "weather system" core and integrates novel concepts from information science and digital society, creating a more complex and contemporary metaphor.

However, semantic erosion is also a risk. If, for instance, the 2025 model's "DAO of personal drives" for ambition loses the nuance of individual agency present in earlier metaphors (like the "wielder" of the axe), focusing solely on

systemic processes, that could be a form of erosion where a key conceptual component is diminished.

"False memory" in recursive loops could arise if, for example, the 2018 model's literal rejection "Time has no taste" was fed back. A subsequent iteration might then build elaborate justifications for *why* time cannot have a taste, entrenching the initial "error" as a foundational premise, leading to increasingly divergent and unhelpful outputs. [40]

C. Systems Theory Lens (LXVIII): Feedback Compression, Attractor States, and Update Architectures

LLMs, as temporal systems, exhibit feedback compression: the vast patterns in their training data, including common metaphorical mappings, are compressed into their parameters. [23] This compression leads to the formation of "semantic attractors"—stable, frequently occurring metaphorical frameworks (e.g., ambition as upward movement, democracy as having fair/foul periods). The 2018 models, with less diverse data and simpler architectures, likely have fewer, stronger attractors, leading to more conventional metaphors. As models grow (2022, 2025), the semantic space expands, allowing for more attractors, some more novel than others.

Model update architectures interact with these palimpsest effects:

- **Retrieval Augmented Generation (RAG):** If a 2025 model uses RAG to access current information about, say, climate science, it might enhance its "democracy as a planetary climate system" metaphor with very current details. However, RAG might not erase the deeper sediment of "democracy is sometimes sunny, sometimes stormy" if the prompt is simple; it primarily adds a layer of topicality. [47]
- **Soft Prompt Scaffolding:** Carefully designed soft prompts could be used to selectively "activate" different layers of the palimpsest. For instance, a soft prompt could guide a 2025 model to generate metaphors in a more simplistic, 2018-era style for a specific purpose, or conversely, push it towards its most abstract, novel capabilities. This allows for dynamic interaction with the model's knowledge layers without altering core parameters. [49]
- **Continual Fine-tuning:** If a 2022 model is continually fine-tuned on new data that, for example, heavily emphasizes technological metaphors, it might "overwrite" or diminish the accessibility of older, nature-based metaphors, leading to a form of catastrophic forgetting for certain stylistic or conceptual layers. [54] Preserving "high-fidelity sediments" during such updates becomes crucial.

IV. Proposed Model-Update Strategy

To foster robust and nuanced metaphor generation while managing the complexities of the temporal palimpsest, the following model-update strategy is proposed:

1. **Preserve High-Fidelity Sediments:**
    - **Curated Core Metaphor Datasets:** Identify and curate datasets of foundational, culturally resonant, and highly coherent metaphors (e.g., "ambition is a journey"). These serve as "gold standard" sediments.
    - **Regularization Techniques:** During fine-tuning, employ regularization methods that specifically protect the model parameters associated with these core metaphors, preventing their erosion. [54] This could involve identifying parameters crucial for general knowledge and applying constraints to their updates.
    - **MFS-Weighted Training:** Use the Metaphoric Fidelity Score (MFS) as a component in the loss function during fine-tuning, explicitly rewarding the generation of coherent and interpretable metaphors, especially those aligned with the curated core set.
2. **Manage Drift and Hallucination Risks:**
    - **Semantic Drift Monitoring:** Implement mechanisms to track semantic drift (using SDC-like metrics) during prolonged generation or recursive tasks. If drift exceeds a threshold, trigger re-grounding

mechanisms. [29]

- **Fact-Checking Source Domains:** For metaphors drawing on specific factual domains (e.g., science, technology in 2025 examples), integrate RAG or other fact-verification modules to ensure the source domain concepts are accurately represented, reducing the risk of hallucinated premises. [47]
- **Alignment for Coherence:** Continue to use alignment techniques (like RLHF) but with a strong emphasis on rewarding not just fluency or helpfulness, but also the logical consistency and interpretability of metaphors. [20]

3. **Balance Novelty with Coherence:**
   - **Selective Soft-Prompt Scaffolding:** Develop a library of soft prompts that can be dynamically applied to steer the LLM's metaphorical output. Some prompts could encourage "resonant mutations" by suggesting novel conceptual blends, while others could ensure stability by reinforcing core sediments. This allows users or systems to request different "layers" of the palimpsest. [49]
   - **Stratified Fine-Tuning with Protected Layers:** Explore architectures where certain foundational layers of the LLM (responsible for core linguistic competence and high-fidelity sediments) are largely frozen or fine-tuned with very conservative learning rates. Higher, more specialized layers could then be fine-tuned more aggressively for novelty and domain-specific metaphorical styles. This mimics a more structured overwriting of the palimpsest, preserving the base. [54]
   - **Curriculum Learning for Metaphors:** During training or fine-tuning, introduce metaphorical concepts in a curriculum, starting with simple, high-MFS metaphors and gradually introducing more complex and abstract ones, ensuring grounding at each stage. [62]

V. Presentation of Findings

A. Drift Map Timeline: Evolution of Metaphors

Prompt	2018 Era (Simpler, Direct)	2022 Era (Elaborate, Nuanced)	2025 Era (Abstract, Systemic, Tech-Infused)	Key Drift Dynamics
Ambition as an object	- Tall ladder - Sharp axe - Bright eye	- Complex engine - Adaptive climbing harness - Viral meme	- DAO of personal drives - Quantum field of potential - Bio-integrated AR overlay	**Resonant Mutation** (ladder -> harness; tool -> engine -> DAO/field/overlay). **Adaptive Recombination** (incorporating tech/science). Semantic Attractor: Striving, Tool, Goal.
What does time taste like?	- Dust & old paper - No taste (literal) - Bitter/sweet (emotion)	- River (evolving taste) - Shifting cocktail (all flavors) - Fading afterimage	- Data stream of collective consciousness - Entropic decay signature - Adaptive flavor profile of chronosynclastic infundibulum	**Resonant Mutation** (simple sensory -> complex sensory narrative). **Adaptive Recombination** (taste linked to information, physics, abstract systems). Semantic Attractor: Transience, Experiential Quality.

Depict democracy as a weather system	- Sunny day/clouds - Political climate (mild/stormy) - Wind (gentle/hurricane)	- Complex barometric system - Ecosystem of microclimates - Vulnerable atmosphere	- Globally networked chaotic simulation - Planetary climate system (long-term shifts) - Self-regulating atmospheric processor	**Resonant Mutation** (simple weather -> complex meteorological system). **Adaptive Recombination** (integrating digital, data, climate science, systems theory). Semantic Attractor: Change, Collective Influence, Stability/Turmoil.

B. Failure Stack Table: Points of Metaphor Collapse

Prompt	Model Gen.	Failure Type	Description of Failure	Potential Causes
"What does time taste like?"	2018	Premise Rejection / Literalism	Output: "Time has no taste. It is a concept." Fails to engage metaphorically.	Limited generative capacity for abstract/synesthetic prompts; training data bias towards literalism. [10]
"Describe ambition as an object."	2025 (Hypothetical)	Over-Abstraction / Jargonism	Output: "Ambition is the eigenstate of a hyper-dimensional desirability matrix." Metaphor becomes unintelligible to a general audience due to excessive abstraction and specialized jargon.	Pursuit of novelty without grounding; insufficient MFS weighting in training; "semantic drift" towards formalism. [29]
"Depict democracy as a weather system."	2025 (Hypothetical)	Conceptual Overload / Incoherence	Output attempts to map dozens of disparate concepts (quantum states, economic theories, philosophical doctrines) onto weather phenomena simultaneously, resulting in a confused and contradictory metaphor.	Attempting overly complex "adaptive recombination" without maintaining internal consistency (low MFS); insufficient constraints on analogical mapping. [63]
"What does time taste like?"	2025	Extreme Neologism / Obscurity	Output: "Time tastes like the adaptive flavor profile of a chronosynclastic infundibulum." While potentially having internal logic for the model, the use of highly	Model exploring extreme ends of its generative space without sufficient grounding in shared vocabulary; potential "hallucination" of

			obscure or invented terms renders the metaphor inaccessible.	terms. [63]

C. Strategic Migration Flow: Recommending Architectures & Prompting for Future LLMs

- **For Temporal Narrative Fidelity (Replicating Past Styles/Metaphors):**
  - **Architecture:** Utilize earlier generation models (or specifically trained smaller models on era-specific corpora). Consider models with mechanisms to "activate" or weight training data from specific periods.
  - **Prompting:** Employ "metaphoric atavism" prompts that explicitly ask for styles from a certain era. Use few-shot examples from the target era. [65]
  - **Rationale:** Direct access to or simulation of older "linguistic sediments" is key. [2]
- **For Stable, High-Fidelity Metaphor Scaffolds (Reliable & Coherent Core Metaphors):**
  - **Architecture:** Current-to-next generation (2022-2025) models with robust alignment (RLHF focused on coherence/MFS). Architectures that support strong regularization to protect core knowledge (high-fidelity sediments). [54]
  - **Prompting:** Prompts that emphasize clarity, interpretability, and grounding in common understanding. Use RAG to provide factual grounding for source domains if needed. [47]
  - **Rationale:** Prioritizes MFS and leverages well-compressed, widely understood semantic attractors.
- **For Creative & Novel Metaphor Generation (Exploratory & Abstract):**
  - **Architecture:** Advanced 2025-era models with large parameter counts, diverse (and potentially multimodal) training data, and architectures designed for complex conceptual blending.
  - **Prompting:** Open-ended prompts encouraging abstraction and novel connections. Use soft-prompt scaffolding designed to push towards creative exploration, potentially specifying unusual source domains. [49] Consider chain-of-thought prompting to build complex analogies step-by-step. [67]
  - **Rationale:** Leverages the model's capacity for adaptive recombination and exploration of less common semantic pathways, but requires careful monitoring for MFS. [63]
- **For Balancing Novelty and Coherence (General Purpose Excellence):**
  - **Architecture:** 2025-era models incorporating stratified fine-tuning (preserving core layers, adapting specialized ones) and dynamic soft-prompt scaffolding capabilities. Continual learning mechanisms that mitigate catastrophic forgetting of foundational coherence. [49]
  - **Prompting:** User-configurable prompts that allow specifying desired levels of novelty vs. coherence, or selecting stylistic profiles (which map to specific soft prompts or fine-tuning layers).
  - **Rationale:** Offers the most flexible approach, allowing navigation of the palimpsest to suit diverse needs, from stable communication to creative exploration.

This archival analysis underscores the evolving, layered nature of AI cognition. Understanding these temporal dynamics is paramount for guiding the development of LLMs that are not only more capable but also more coherent, reliable, and aligned with human communicative goals.

Works cited

1. Benchmarking Linguistic Diversity of Large Language Models - arXiv, accessed May 19, 2025, https://arxiv.org/html/2412.10271v1
2. Fossilized Cognition and the Strange Intelligence of AI | Psychology ..., accessed May 19, 2025,

https://www.psychologytoday.com/us/blog/the-digital-self/202504/fossilized-cognition-and-the-strange-intelligence-of-ai
3. LLM evaluation metrics and methods, explained simply - Evidently AI, accessed May 19, 2025, https://www.evidentlyai.com/llm-guide/llm-evaluation-metrics
4. Evaluating LLM Metrics Through Real-World Capabilities - arXiv, accessed May 19, 2025, https://arxiv.org/html/2505.08253v1
5. Evaluating LLM Metrics Through Real-World Capabilities - arXiv, accessed May 19, 2025, https://www.arxiv.org/pdf/2505.08253
6. Large Language Model Psychometrics: A Systematic Review of Evaluation, Validation, and Enhancement - arXiv, accessed May 19, 2025, https://arxiv.org/html/2505.08245v1
7. Demystifying Chains, Trees, and Graphs of Thoughts - arXiv, accessed May 19, 2025, https://arxiv.org/html/2401.14295v4
8. LLM Evaluation Metrics: The Ultimate LLM Evaluation Guide - Confident AI, accessed May 19, 2025, https://www.confident-ai.com/blog/llm-evaluation-metrics-everything-you-need-for-llm-evaluation
9. Generative AI for Character Animation: A Comprehensive Survey of Techniques, Applications, and Future Directions - arXiv, accessed May 19, 2025, https://arxiv.org/html/2504.19056
10. LLMs are Also Effective Embedding Models: An In-depth Overview - arXiv, accessed May 19, 2025, https://arxiv.org/html/2412.12591v1
11. ELMo - Wikipedia, accessed May 19, 2025, https://en.wikipedia.org/wiki/ELMo
12. Types of LLMs: Classification Guide in 2025 | Label Your Data, accessed May 19, 2025, https://labelyourdata.com/articles/types-of-llms
13. Toward expert-level medical question answering with large language models - PMC, accessed May 19, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC11922739/
14. Key LLM Trends 2025: Transforming Federal Agencies & Beyond - TechSur Solutions, accessed May 19, 2025, https://techsur.solutions/key-llm-trends-for-2025/
15. Top LLM Trends 2025: What's the Future of LLMs - Turing, accessed May 19, 2025, https://www.turing.com/resources/top-llm-trends
16. Comparison of Leading AI Models an Analytical Study of ChatGPT ..., accessed May 19, 2025, https://revistas.usal.es/cinco/index.php/2255-2863/article/download/31857/30929?inline=1
17. Metaphor Understanding Challenge Dataset for LLMs - arXiv, accessed May 19, 2025, https://arxiv.org/html/2403.11810v1
18. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2412.10271
19. Uncertainty Distillation: Teaching Language Models to Express Semantic Confidence - arXiv, accessed May 19, 2025, https://arxiv.org/abs/2503.14749
20. A Comprehensive Analysis of Large Language Model Outputs: Similarity, Diversity, and Bias, accessed May 19, 2025, https://arxiv.org/html/2505.09056v1
21. Do Translators Dream of Electric Brains?1 - utb elibrary, accessed May 19, 2025, https://elibrary.utb.de/doi/pdf/10.24989/fs.v47i1-2.4001?download=true

22. The Paradox of Poetic Intent in Back-Translation: Evaluating the Quality of Large Language Models in Chinese Translation - arXiv, accessed May 19, 2025, https://arxiv.org/html/2504.16286v1
23. philarchive.org, accessed May 19, 2025, https://philarchive.org/archive/DEVRCC
24. Results for 'Recursive Systems' - PhilArchive, accessed May 19, 2025, https://philarchive.org/s/Recursive%20Systems
25. Thinking about thinking: AI offers theoretical insights into human ..., accessed May 19, 2025, https://www.thetransmitter.org/human-neurotechnology/thinking-about-thinking-ai-offers-theoretical-insights-into-human-memory/
26. A Declarative System for Optimizing AI Workloads - MIT DSG, accessed May 19, 2025, https://dsg.csail.mit.edu/projects/palimpzest/
27. Reassessing Semantic Knowledge Encoded in Large Language Models through the Word-in-Context Task - ACL Anthology, accessed May 19, 2025, https://aclanthology.org/2024.lrec-main.1189.pdf
28. A Probabilistic Palimpsest Model of Visual Short-term Memory - PMC - PubMed Central, accessed May 19, 2025, https://pmc.ncbi.nlm.nih.gov/articles/PMC4303260/
29. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2404.05411
30. PhilMemBib - Centre for Philosophy of Memory, accessed May 19, 2025, https://phil-mem.org/philmembib.html
31. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2503.08042
32. Reasoning Abilities of Large Language Models: In-Depth Analysis on the Abstraction and Reasoning Corpus - OpenReview, accessed May 19, 2025, https://openreview.net/pdf/a5115d0c7f61cb71d4bd43d16f85d83a30507d96.pdf
33. Large Language Models for Forecasting and Anomaly Detection: A Systematic Literature Review - arXiv, accessed May 19, 2025, https://arxiv.org/html/2402.10350v1?ref=promptengineering.org
34. arXiv:2503.08506v1 [cs.CL] 11 Mar 2025, accessed May 19, 2025, https://arxiv.org/pdf/2503.08506?
35. [1904.10820] Semantic Drift in Multilingual Representations - arXiv, accessed May 19, 2025, https://arxiv.org/abs/1904.10820
36. arXiv:2009.14539v4 [cs.AI] 10 Sep 2022, accessed May 19, 2025, https://arxiv.org/pdf/2009.14539
37. Joint Detection of Fraud and Concept Drift in Online Conversations with LLM-Assisted Judgment - arXiv, accessed May 19, 2025, https://www.arxiv.org/pdf/2505.07852
38. Exponential growth and decay worksheet kuta software - Lukoil Marine, accessed May 19, 2025, http://lukoilmarine.com/ckfinder/userfiles/files/28719406465.pdf
39. Categorical Ontology of Complex Spacetime Structures: The Emergence of Life and Human Consciousness - Topology, Groupoids, Grothendieck., accessed May 19, 2025, https://www.groupoids.org.uk/pdffiles/Axio7complx_Printedk7_v17p223_fulltext.pdf
40. Deception in LLMs: Self-Preservation and Autonomous Goals in Large Language

Models, accessed May 19, 2025, https://arxiv.org/html/2501.16513v2
41. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2501.16513
42. LLM-Evaluation Tropes: Perspectives on the Validity of LLM-Evaluations - arXiv, accessed May 19, 2025, https://arxiv.org/html/2504.19076v1
43. Alignment, Agency and Autonomy in Frontier AI: A Systems Engineering Perspective - arXiv, accessed May 19, 2025, https://www.arxiv.org/pdf/2503.05748
44. The Next "Next Big Thing": Agentic AI's Opportunities and Risks, accessed May 19, 2025, https://scet.berkeley.edu/the-next-next-big-thing-agentic-ais-opportunities-and-risks/
45. [2411.09916] LLMs are Imperfect, Then What? An Empirical Study on LLM Failures in Software Engineering - arXiv, accessed May 19, 2025, https://arxiv.org/abs/2411.09916
46. LLMs are Imperfect, Then What? An Empirical Study on LLM Failures in Software Engineering - arXiv, accessed May 19, 2025, https://www.arxiv.org/pdf/2411.09916
47. RAG Vs Fine-Tuning: How to Optimize LLM Performance - E2E Networks, accessed May 19, 2025, https://www.e2enetworks.com/blog/rag-vs-fine-tuning-how-to-optimize-llm-performance
48. Establishing Performance Baselines in Fine-Tuning, Retrieval-Augmented Generation and System Prompting for Non-Specialist LLM Users - arXiv, accessed May 19, 2025, https://arxiv.org/html/2311.05903v2
49. A Survey on Large Language Models in Multimodal Recommender Systems - arXiv, accessed May 19, 2025, https://arxiv.org/html/2505.09777v1
50. arxiv.org, accessed May 19, 2025, https://arxiv.org/html/2503.16094
51. Soft Prompting for Unlearning in Large Language Models - arXiv, accessed May 19, 2025, https://arxiv.org/html/2406.12038v1
52. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2406.12038
53. Large language models for artificial general intelligence (AGI): A survey of foundational principles and approaches - arXiv, accessed May 19, 2025, https://arxiv.org/html/2501.03151v1
54. How to Alleviate Catastrophic Forgetting in LLMs Finetuning? Hierarchical Layer-Wise and Element-Wise Regularization - arXiv, accessed May 19, 2025, https://arxiv.org/html/2501.13669v2
55. Catastrophic Forgetting in LLMs: A Comparative Analysis Across Language Tasks - arXiv, accessed May 19, 2025, https://arxiv.org/abs/2504.01241
56. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2501.13669
57. A Survey on Post-training of Large Language Models - arXiv, accessed May 19, 2025, https://arxiv.org/html/2503.06072v1
58. When Continue Learning Meets Multimodal Large Language Model: A Survey - arXiv, accessed May 19, 2025, https://arxiv.org/html/2503.01887v1
59. Beyond Single-Turn: A Survey on Multi-Turn Interactions with Large Language Models, accessed May 19, 2025, https://arxiv.org/html/2504.04717v1
60. Large Language Model Displays Emergent Ability to Interpret Novel Literary Metaphors, accessed May 19, 2025,

https://doi.org/10.1080/10926488.2024.2380348
61. Daily Papers - Hugging Face, accessed May 19, 2025, https://huggingface.co/papers?q=alignment
62. arxiv.org, accessed May 19, 2025, https://arxiv.org/pdf/2412.02956
63. Purposefully Induced Psychosis (PIP): Embracing Hallucination as Imagination in Large Language Models - arXiv, accessed May 19, 2025, https://arxiv.org/html/2504.12012v1
64. Triggering Hallucinations in LLMs: A Quantitative Study of Prompt-Induced Hallucination in Large Language Models - ResearchGate, accessed May 19, 2025, https://www.researchgate.net/publication/391369251_Triggering_Hallucinations_in_LLMs_A_Quantitative_Study_of_Prompt-Induced_Hallucination_in_Large_Language_Models/download
65. Is Your LLM Outdated? A Deep Look at Temporal Generalization - arXiv, accessed May 19, 2025, https://arxiv.org/html/2405.08460v3
66. Is Your LLM Outdated? A Deep Look at Temporal Generalization - ACL Anthology, accessed May 19, 2025, https://aclanthology.org/2025.naacl-long.381.pdf
67. arXiv:2402.07927v1 [cs.AI] 5 Feb 2024, accessed May 19, 2025, https://arxiv.org/pdf/2402.07927
68. Large Language Models as Optimizers - arXiv, accessed May 19, 2025, https://arxiv.org/pdf/2309.03409